

## chi\_sq\_30 7030 split .05 threshold

January 2, 2023

```
[ ]: # Importing the packages
import sys
import numpy as np
np.set_printoptions(threshold=sys.maxsize)
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
import sklearn
import random
from sklearn.metrics import
    ↪confusion_matrix, accuracy_score, classification_report, RocCurveDisplay, ConfusionMatrixDisplay

[ ]: pd.set_option('display.max_rows', None)
pd.set_option('display.max_columns', None)
pd.set_option('display.width', None)
pd.set_option('display.max_colwidth', None)

[ ]: # Importing the dataset
df = pd.read_csv('dataset_30.csv')
df.drop(['index'], axis=1, inplace=True)
#df.head()

[ ]: # if your dataset contains missing value, check which column has missing values
#df.isnull().sum()

[ ]: #df.dropna(inplace=True)

[ ]: from sklearn import preprocessing

col = df.columns[:]

lab_en= preprocessing.LabelEncoder()

for c in col:
    df[c]= lab_en.fit_transform(df[c])

#df.head(50)
```

```
[ ]: a=len(df[df.Result==0])
      b=len(df[df.Result==1])
```

```
[ ]: print("Count of Legitimate Websites = ", a)
      print("Count of Phishy Websites = ", b)
```

Count of Legitimate Websites = 4898  
Count of Phishy Websites = 6157

```
[ ]: X = df.drop(['Result'], axis=1, inplace=False)
      #X.head()
      #same work
      ##inplace true modifies the og data & does not return anything
      ##inplace false does not modify og data but returns something which we store in
      ↪ a var
      # X= df.drop(columns='Result')
      # X.head()
```

```
[ ]: #df.head()
```

```
[ ]: y = df['Result']
      y = pd.DataFrame(y)
      y.head()
```

```
[ ]:      Result
      0      0
      1      0
      2      0
      3      0
      4      1
```

```
[ ]: # separate dataset into train and test
      from cProfile import label
      from sklearn.model_selection import train_test_split
      X_train, X_test, y_train, y_test = train_test_split(
          X,
          y,
          test_size=0.3,
          random_state=10)

      X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
[ ]: ((7738, 30), (3317, 30), (7738, 1), (3317, 1))
```

```
[ ]: #perform chi square test
      from sklearn.feature_selection import chi2
      f_p_values = chi2(X_train,y_train)
```

```
[ ]: f_p_values
```

```
[ ]: (array([2.24728145e+01, 5.38784812e+01, 5.03249360e+00, 3.83671222e+00,
            1.68176772e+00, 8.20768441e+02, 4.52503451e+02, 2.66189096e+03,
            2.70267968e+02, 7.35475144e-03, 1.55472779e+00, 2.46950144e+00,
            2.10635122e+02, 2.09920590e+03, 3.13269211e+02, 5.21473196e+02,
            5.20045979e-01, 4.51736930e+00, 3.28650457e+00, 1.70791844e+00,
            9.24087420e-02, 8.74179169e-04, 3.80728975e-04, 5.46224163e+01,
            1.36738388e+01, 4.87492812e+02, 6.43890373e+01, 1.65423051e+01,
            1.36921170e+00, 6.09019636e+00]),
      array([2.13138824e-006, 2.13280861e-013, 2.48760576e-002, 5.01417529e-002,
            1.94689729e-001, 1.64702370e-180, 2.05730537e-100, 0.00000000e+000,
            9.91989265e-061, 9.31657326e-001, 2.12438845e-001, 1.16074737e-001,
            9.98392056e-048, 0.00000000e+000, 4.23702310e-070, 2.02315605e-115,
            4.70822069e-001, 3.35523845e-002, 6.98515795e-002, 1.91255666e-001,
            7.61136986e-001, 9.76412766e-001, 9.84432443e-001, 1.46060070e-013,
            2.17462929e-004, 5.00438319e-108, 1.02124797e-015, 4.75766558e-005,
            2.41947370e-001, 1.35933979e-002]))
```

```
[ ]: #The less the p_values the more important that feature is
p_values = pd.Series(f_p_values[1])
p_values.index = X_train.columns
p_values
```

```
[ ]: having_IPhaving_IP_Address      2.131388e-06
      URLURL_Length                  2.132809e-13
      Shortining_Service              2.487606e-02
      having_At_Symbol                5.014175e-02
      double_slash_redirecting        1.946897e-01
      Prefix_Suffix                   1.647024e-180
      having_Sub_Domain               2.057305e-100
      SSLfinal_State                  0.000000e+00
      Domain_registration_length      9.919893e-61
      Favicon                         9.316573e-01
      port                            2.124388e-01
      HTTPS_token                     1.160747e-01
      Request_URL                     9.983921e-48
      URL_of_Anchor                   0.000000e+00
      Links_in_tags                   4.237023e-70
      SFH                             2.023156e-115
      Submitting_to_email             4.708221e-01
      Abnormal_URL                    3.355238e-02
      Redirect                        6.985158e-02
      on_mouseover                    1.912557e-01
      RightClick                      7.611370e-01
      popUpWidnow                     9.764128e-01
      Iframe                          9.844324e-01
```

```

age_of_domain          1.460601e-13
DNSRecord              2.174629e-04
web_traffic            5.004383e-108
Page_Rank              1.021248e-15
Google_Index          4.757666e-05
Links_pointing_to_page 2.419474e-01
Statistical_report     1.359340e-02
dtype: float64

```

```

[ ]: #sort p_values to check which feature has the lowest values
p_values = p_values.sort_values(ascending = False)
p_values

```

```

[ ]: Iframe          9.844324e-01
popUpWidnow         9.764128e-01
Favicon             9.316573e-01
RightClick          7.611370e-01
Submitting_to_email 4.708221e-01
Links_pointing_to_page 2.419474e-01
port                2.124388e-01
double_slash_redirecting 1.946897e-01
on_mouseover        1.912557e-01
HTTPS_token         1.160747e-01
Redirect            6.985158e-02
having_At_Symbol    5.014175e-02
Abnormal_URL        3.355238e-02
Shortining_Service  2.487606e-02
Statistical_report  1.359340e-02
DNSRecord           2.174629e-04
Google_Index        4.757666e-05
having_IPhaving_IP_Address 2.131388e-06
URLURL_Length       2.132809e-13
age_of_domain       1.460601e-13
Page_Rank           1.021248e-15
Request_URL         9.983921e-48
Domain_registration_length 9.919893e-61
Links_in_tags       4.237023e-70
having_Sub_Domain   2.057305e-100
web_traffic         5.004383e-108
SFH                 2.023156e-115
Prefix_Suffix       1.647024e-180
URL_of_Anchor       0.000000e+00
SSLfinal_State      0.000000e+00
dtype: float64

```

```

[ ]: def DropFeature (p_values, threshold):
      drop_feature = set()

```

```

    for index, values in p_values.items():
        if values > threshold or np.isnan(values):
            drop_feature.add(index)
    return drop_feature

```

```

[ ]: drop_feature = DropFeature(p_values, .05)
len(set(drop_feature))

```

```

[ ]: 12

```

```

[ ]: drop_feature

```

```

[ ]: {'Favicon',
      'HTTPS_token',
      'Iframe',
      'Links_pointing_to_page',
      'Redirect',
      'RightClick',
      'Submitting_to_email',
      'double_slash_redirecting',
      'having_At_Symbol',
      'on_mouseover',
      'popUpWidnow',
      'port'}

```

```

[ ]: X_train.drop(drop_feature, axis=1, inplace=True)
X_test.drop(drop_feature, axis=1, inplace=True)

```

```

[ ]: len(df.columns)

```

```

[ ]: 31

```

```

[ ]: print("Training set has {} samples.".format(X_train.shape[0]))
print("Testing set has {} samples.".format(X_test.shape[0]))

```

Training set has 7738 samples.

Testing set has 3317 samples.

```

[ ]: from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression

# defining parameter range
param_grid = {'penalty' : ['l2'],
              'C' : [0.1, 1, 10, 20, 30],
              'solver' : ['lbfgs', 'newton-cg', 'liblinear', 'sag', 'saga'],
              'max_iter' : [2500, 5000]}

```

```

grid_logr = GridSearchCV(LogisticRegression(), param_grid, refit = True, cv =
↳10, verbose = 3, n_jobs = -1)

# fitting the model for grid search
grid_logr.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_logr.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_logr.best_estimator_)
print(grid_logr.best_score_)

```

Fitting 10 folds for each of 50 candidates, totalling 500 fits  
{'C': 1, 'max\_iter': 2500, 'penalty': 'l2', 'solver': 'sag'}  
LogisticRegression(C=1, max\_iter=2500, solver='sag')  
0.9246594529184259

```

[ ]: logr_model = grid_logr.best_estimator_

# Performing training
#logr_model = logr.fit(X_train, y_train.values.ravel())

```

```

[ ]: logr_predict = logr_model.predict(X_test)

```

```

[ ]: # from sklearn.metrics import confusion_matrix, accuracy_score
# cm = confusion_matrix(y_test, dct_pred)
# ac = accuracy_score(y_test, dct_pred)

```

```

[ ]: print ("Accuracy of logr classifier : ", accuracy_score(y_test,
↳logr_predict)*100)

```

Accuracy of logr classifier : 91.98070545673802

```

[ ]: print(classification_report(y_test, logr_predict))

```


	precision	recall	f1-score	support
0	0.92	0.89	0.91	1444
1	0.92	0.94	0.93	1873
accuracy			0.92	3317
macro avg	0.92	0.92	0.92	3317
weighted avg	0.92	0.92	0.92	3317

```

[ ]: sns.heatmap(confusion_matrix(y_test, logr_predict), annot=True, fmt='g',
↳cmap='Blues')
plt.title("LogisticRegression")

```

```
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
```



```
[ ]: # from sklearn.neighbors import KNeighborsClassifier

# #training_accuracy=[]
# test_accuracy=[]

# neighbors=range(1,10)
# ##values.ravel() converts vector y to flattened array
# for i in neighbors:
#     knn=KNeighborsClassifier(n_neighbors=i)
#     knn_model = knn.fit(X_train,y_train.values.ravel())
#     #training_accuracy.append(knn.score(X_train,y_train.values.ravel()))
#     test_accuracy.append(knn_model.score(X_test,y_test.values.ravel()))
```

```
[ ]: # plt.plot(neighbors,test_accuracy,label="test accuracy")
# plt.ylabel("Accuracy")
```

```
# plt.xlabel("number of neighbors")
# plt.legend()
# plt.show()
```

```
[ ]: from sklearn.neighbors import KNeighborsClassifier

# defining parameter range
param_grid = {'n_neighbors': [1,2,3,4,5,6,7,8,9,10]}

grid_knn = GridSearchCV(KNeighborsClassifier(), param_grid, refit = True, cv = 10,
↳ verbose = 3, n_jobs = -1)

# fitting the model for grid search
grid_knn.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_knn.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_knn.best_estimator_)
print(grid_knn.best_score_)
```

```
Fitting 10 folds for each of 10 candidates, totalling 100 fits
{'n_neighbors': 1}
KNeighborsClassifier(n_neighbors=1)
0.9519234433446654
```

```
[ ]: knn_model = grid_knn.best_estimator_
#knn_model = knn.fit(X_train,y_train.values.ravel())
```

```
[ ]: #print ("Accuracy of knn classifier: ", max(test_accuracy)*100)
knn_predict = knn_model.predict(X_test)
```

```
[ ]: print('The accuracy of knn Classifier is: ', 100.0 * accuracy_score(y_test,
↳ knn_predict))
```


```
The accuracy of knn Classifier is: 95.29695507989146
```

```
[ ]: print(classification_report(y_test, knn_predict))
```

	precision	recall	f1-score	support
0	0.95	0.94	0.95	1444
1	0.95	0.96	0.96	1873
accuracy			0.95	3317
macro avg	0.95	0.95	0.95	3317
weighted avg	0.95	0.95	0.95	3317



```
[ ]: sns.heatmap(confusion_matrix(y_test, knn_predict), annot=True, fmt='g',  
    cmap='Blues')  
plt.title("KNeighborsClassifier")  
plt.xlabel('Predicted Class')  
plt.ylabel('Original Class')  
plt.show()
```



```
[ ]: # # here is the change  
# knn_y_pred_proba = knn.predict_proba(X_test)  
# knn_y_pred_proba_positive = knn_y_pred_proba[:, 1]  
  
# RocCurveDisplay.from_predictions(y_test,knn_y_pred_proba_positive)  
  
# fig, ax = plt.subplots()  
# RocCurveDisplay.from_estimator(  
#     logreg, X_test, y_test, ax = ax)  
  
# logreg_y_decision = logreg.decision_function(X_test)
```

```
# metrics.RocCurveDisplay.
↪from_predictions(y_test, logreg_y_decision, ax=ax, name="logreg predictions")
```

```
[ ]: from sklearn.svm import SVC

# defining parameter range
param_grid = {'C': [0.1, 1, 10],
              'gamma': [1, 0.1, 0.01],
              'kernel': ['linear', 'poly', 'rbf', 'sigmoid']}

grid_svc = GridSearchCV(SVC(), param_grid, refit = True, cv = 10, verbose = 3, ↪
↪n_jobs = -1)

# fitting the model for grid search
grid_svc.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_svc.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_svc.best_estimator_)
print(grid_svc.best_score_)
```

Fitting 10 folds for each of 36 candidates, totalling 360 fits  
{'C': 1, 'gamma': 1, 'kernel': 'rbf'}  
SVC(C=1, gamma=1)  
0.9564464100069865

```
[ ]: svc_model = grid_svc.best_estimator_
#svc_model = svc.fit(X_train, y_train.values.ravel())
```

```
[ ]: svc_predict = svc_model.predict(X_test)
```


```
[ ]: print('The accuracy of svc Classifier is: ', 100.0 * accuracy_score(y_test, ↪
↪svc_predict))
```

The accuracy of svc Classifier is: 96.11094362375641

```
[ ]: print(classification_report(y_test, svc_predict))
```

	precision	recall	f1-score	support
0	0.96	0.95	0.95	1444
1	0.96	0.97	0.97	1873
accuracy			0.96	3317
macro avg	0.96	0.96	0.96	3317
weighted avg	0.96	0.96	0.96	3317

```
[ ]: sns.heatmap(confusion_matrix(y_test, svc_predict), annot=True, fmt='g',  
    cmap='Blues')  
plt.title("SVC")  
plt.xlabel('Predicted Class')  
plt.ylabel('Original Class')  
plt.show()
```



```
[ ]: from sklearn.svm import NuSVC  
  
# defining parameter range  
param_grid = {'nu': [0.1, 0.5],  
              'gamma': [1, 0.1, 0.01],  
              'kernel': ['linear', 'poly', 'rbf', 'sigmoid']}  
  
grid_nusvc = GridSearchCV(NuSVC(), param_grid, refit = True, verbose = 3, cv =  
    10, n_jobs = -1)  
  
# fitting the model for grid search  
grid_nusvc.fit(X_train, y_train.values.ravel())
```

```
# print best parameter after tuning
print(grid_nusvc.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_nusvc.best_estimator_)
print(grid_nusvc.best_score_)
```

```
Fitting 10 folds for each of 24 candidates, totalling 240 fits
{'gamma': 1, 'kernel': 'rbf', 'nu': 0.1}
NuSVC(gamma=1, nu=0.1)
0.9547658206056472
```

```
[ ]: nusvc_model = grid_nusvc.best_estimator_
      #nusvc_model = nusvc.fit(X_train, y_train.values.ravel())
```

```
[ ]: nusvc_predict = nusvc_model.predict(X_test)
```


```
[ ]: print('The accuracy of nusvc Classifier is: ', 100.0 * accuracy_score(y_test,
      ↪nusvc_predict))
```

```
The accuracy of nusvc Classifier is: 96.68375037684655
```

```
[ ]: print(classification_report(y_test, nusvc_predict))
```

	precision	recall	f1-score	support
0	0.97	0.95	0.96	1444
1	0.96	0.98	0.97	1873
accuracy			0.97	3317
macro avg	0.97	0.97	0.97	3317
weighted avg	0.97	0.97	0.97	3317

```
[ ]: sns.heatmap(confusion_matrix(y_test, nusvc_predict), annot=True, fmt='g',
      ↪cmap='Blues')
plt.title("NuSVC")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
```



```
[ ]: from sklearn.svm import LinearSVC

# defining parameter range
param_grid = {'C': [0.1, 1, 10, 20, 30],
              'penalty': ['l1', 'l2'],
              'loss': ['squared_hinge'],
              'dual': [False],
              'tol': [.1, .01, .001]}

grid_lsvc = GridSearchCV(LinearSVC(), param_grid, refit = True, verbose = 3, cv=
    ↪ 10, n_jobs = -1)

# fitting the model for grid search
grid_lsvc.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_lsvc.best_params_)

# print how our model looks after hyper-parameter tuning
```

```
print(grid_lsvc.best_estimator_)
print(grid_lsvc.best_score_)
```

Fitting 10 folds for each of 30 candidates, totalling 300 fits  
{'C': 1, 'dual': False, 'loss': 'squared\_hinge', 'penalty': 'l2', 'tol': 0.001}  
LinearSVC(C=1, dual=False, tol=0.001)  
0.9240139595054002

```
[ ]: lsvc_model = grid_lsvc.best_estimator_
      #lsvc_model = lsvc.fit(X_train, y_train.values.ravel())
```

```
[ ]: lsvc_predict = lsvc_model.predict(X_test)
```


```
[ ]: print('The accuracy of lsvc Classifier is: ', 100.0 * accuracy_score(y_test,
      ↪lsvc_predict))
```

The accuracy of lsvc Classifier is: 92.1615917998191

```
[ ]: print(classification_report(y_test, lsvc_predict))
```

	precision	recall	f1-score	support
0	0.93	0.89	0.91	1444
1	0.92	0.95	0.93	1873
accuracy			0.92	3317
macro avg	0.92	0.92	0.92	3317
weighted avg	0.92	0.92	0.92	3317

```
[ ]: sns.heatmap(confusion_matrix(y_test, lsvc_predict), annot=True, fmt='g',
      ↪cmap='Blues')
plt.title("LinearSVC")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
```



```
[ ]: from sklearn.ensemble import AdaBoostClassifier

# defining parameter range
param_grid = {'n_estimators': [40,50,100,200,300]}

grid_ada = GridSearchCV(AdaBoostClassifier(), param_grid, refit = True, verbose=
    ↪ 3, cv = 10, n_jobs = -1)

# fitting the model for grid search
grid_ada.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_ada.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_ada.best_estimator_)
print(grid_ada.best_score_)
```

Fitting 10 folds for each of 5 candidates, totalling 50 fits  
 {'n\_estimators': 200}

```
AdaBoostClassifier(n_estimators=200)
0.9328011606178818
```

```
[ ]: ada_model = grid_ada.best_estimator_  
      #ada_model = ada.fit(X_train,y_train.values.ravel())
```

```
[ ]: ada_predict = ada_model.predict(X_test)
```

```
[ ]: print('The accuracy of Ada Boost Classifier is: ', 100.0 *  
      ↪accuracy_score(ada_predict,y_test))
```


The accuracy of Ada Boost Classifier is: 93.21676213445885

```
[ ]: print(classification_report(y_test, ada_predict))
```

	precision	recall	f1-score	support
0	0.94	0.90	0.92	1444
1	0.93	0.96	0.94	1873
accuracy			0.93	3317
macro avg	0.93	0.93	0.93	3317
weighted avg	0.93	0.93	0.93	3317

```
[ ]: sns.heatmap(confusion_matrix(y_test, ada_predict), annot=True, fmt='g',  
      ↪cmap='Blues')  
plt.title("AdaBoostClassifier")  
plt.xlabel('Predicted Class')  
plt.ylabel('Original Class')  
plt.show()
```





```
[ ]: from xgboost import XGBClassifier

# defining parameter range
param_grid = {
    "gamma": [.01, .1, .5],
    "n_estimators": [50,100,150,200,250]
}

grid_xgb = GridSearchCV(XGBClassifier(), param_grid, refit = True, verbose = 3,
    ↪cv = 10, n_jobs = -1)

# fitting the model for grid search
grid_xgb.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_xgb.best_params_)

# print how our model looks after hyper-parameter tuning
```

```
print(grid_xgb.best_estimator_)
print(grid_xgb.best_score_)
```

Fitting 10 folds for each of 15 candidates, totalling 150 fits

```
{'gamma': 0.01, 'n_estimators': 200}
```

```
XGBClassifier(base_score=0.5, booster='gbtree', callbacks=None,
               colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,
               early_stopping_rounds=None, enable_categorical=False,
               eval_metric=None, gamma=0.01, gpu_id=-1, grow_policy='depthwise',
               importance_type=None, interaction_constraints='',
               learning_rate=0.300000012, max_bin=256, max_cat_to_onehot=4,
               max_delta_step=0, max_depth=6, max_leaves=0, min_child_weight=1,
               missing=nan, monotone_constraints='()', n_estimators=200,
               n_jobs=0, num_parallel_tree=1, predictor='auto', random_state=0,
               reg_alpha=0, reg_lambda=1, ...)
```

0.9609702123676671

```
[ ]: xgb_model = grid_xgb.best_estimator_
      #xgb_model = xgb.fit(X_train,y_train)
```

```
[ ]: xgb_predict=xgb_model.predict(X_test)
```


```
[ ]: print('The accuracy of XGBoost Classifier is: ', 100.0 *
      ↪accuracy_score(xgb_predict,y_test))
```

The accuracy of XGBoost Classifier is: 96.11094362375641

```
[ ]: print(classification_report(y_test, xgb_predict))
```

	precision	recall	f1-score	support
0	0.96	0.95	0.96	1444
1	0.96	0.97	0.97	1873
accuracy			0.96	3317
macro avg	0.96	0.96	0.96	3317
weighted avg	0.96	0.96	0.96	3317

```
[ ]: sns.heatmap(confusion_matrix(y_test, xgb_predict), annot=True, fmt='g',
      ↪cmap='Blues')
plt.title("XGBClassifier")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
```



```
[ ]: from sklearn.ensemble import GradientBoostingClassifier

# defining parameter range
param_grid = {
    "learning_rate": [.1,.5,1],
    "n_estimators": [50,100,150,200,250]
}

grid_gbc = GridSearchCV(GradientBoostingClassifier(), param_grid, refit = True,
    verbose = 3, cv = 10, n_jobs = -1)

# fitting the model for grid search
grid_gbc.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_gbc.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_gbc.best_estimator_)
```

```
print(grid_gbc.best_score_)
```

Fitting 10 folds for each of 15 candidates, totalling 150 fits  
{'learning\_rate': 0.5, 'n\_estimators': 250}  
GradientBoostingClassifier(learning\_rate=0.5, n\_estimators=250)  
0.9578681000564933

```
[ ]: gbc_model = grid_gbc.best_estimator_  
      #gbc_model = gbc.fit(X_train,y_train.values.ravel())  
  
      #clf = GradientBoostingClassifier(n_estimators=100, learning_rate=1.0,  
      #    max_depth=1, random_state=0).fit(X_train, y_train)  
      #clf.score(X_test, y_test)
```

```
[ ]: gbc_predict = gbc_model.predict(X_test)
```


```
[ ]: print('The accuracy of GradientBoost Classifier is: ', 100.0 *  
      ↪accuracy_score(gbc_predict,y_test))
```

The accuracy of GradientBoost Classifier is: 96.05064817606271

```
[ ]: print(classification_report(y_test, gbc_predict))
```

	precision	recall	f1-score	support
0	0.96	0.95	0.95	1444
1	0.96	0.97	0.97	1873
accuracy			0.96	3317
macro avg	0.96	0.96	0.96	3317
weighted avg	0.96	0.96	0.96	3317

```
[ ]: sns.heatmap(confusion_matrix(y_test, gbc_predict), annot=True, fmt='g',  
      ↪cmap='Blues')  
plt.title("GradientBoostingClassifier")  
plt.xlabel('Predicted Class')  
plt.ylabel('Original Class')  
plt.show()
```



```
[ ]: # gbc_model.get_params().keys()
```

```
[ ]: # import inspect
# import sklearn
# import xgboost

# models = [xgboost.XGBClassifier]
# for m in models:
#     hyperparams = inspect.signature(m.__init__)
#     print(hyperparams)
# #or
# xgb_model.get_params().keys()
```

```
[ ]: from sklearn.ensemble import BaggingClassifier
from sklearn.tree import DecisionTreeClassifier

# defining parameter range
param_grid = {
    "base_estimator": [DecisionTreeClassifier()],
    "n_estimators": [50,100,150,200,250]
```

```

}

grid_bag = GridSearchCV(BaggingClassifier(), param_grid, refit = True, verbose_
↳ = 3, cv = 10, n_jobs = -1)

# fitting the model for grid search
grid_bag.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_bag.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_bag.best_estimator_)
print(grid_bag.best_score_)

```

```

Fitting 10 folds for each of 5 candidates, totalling 50 fits
{'base_estimator': DecisionTreeClassifier(), 'n_estimators': 50}
BaggingClassifier(base_estimator=DecisionTreeClassifier(), n_estimators=50)
0.9581269994083256

```

```

[ ]: bag_model = grid_bag.best_estimator_
      #bag_model = bag.fit(X_train, y_train.values.ravel())

```

```

[ ]: bag_predict = bag_model.predict(X_test)

```

```

[ ]: print('The accuracy of Bagging Classifier is: ', 100.0 * _
↳ accuracy_score(y_test, bag_predict))

```

```

The accuracy of Bagging Classifier is: 96.2315345191438

```

```

[ ]: print(classification_report(y_test, bag_predict))

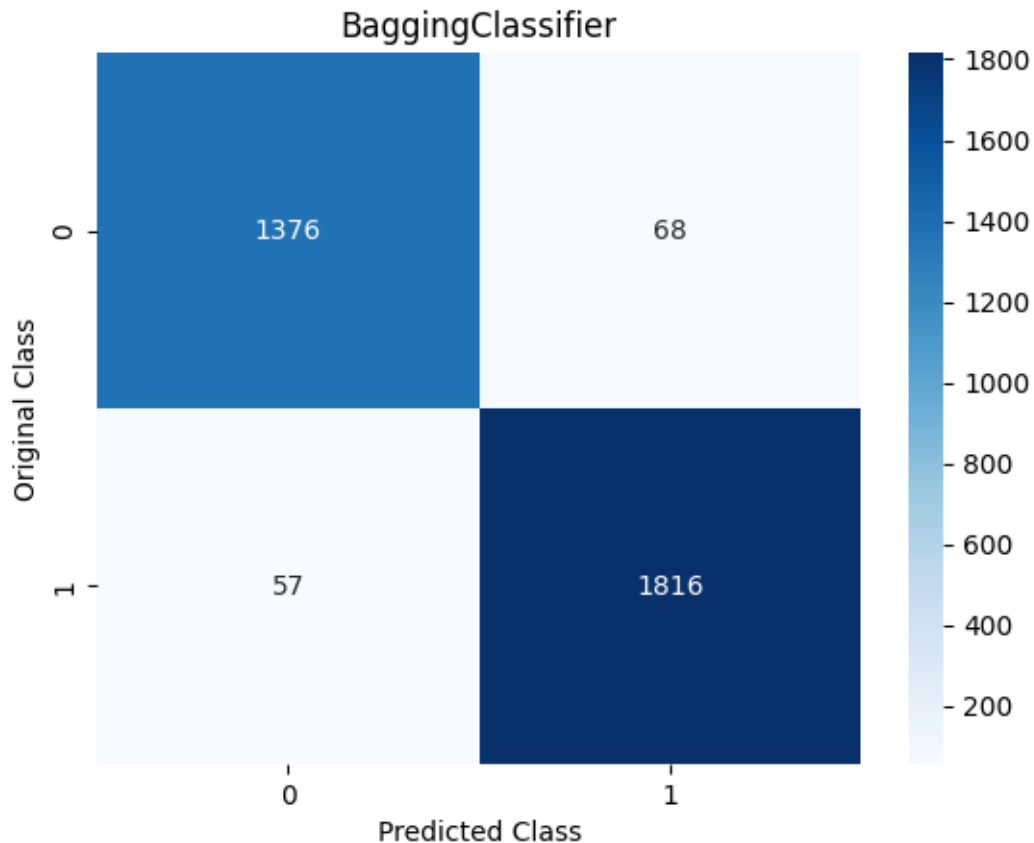
```

	precision	recall	f1-score	support
0	0.96	0.95	0.96	1444
1	0.96	0.97	0.97	1873
accuracy			0.96	3317
macro avg	0.96	0.96	0.96	3317
weighted avg	0.96	0.96	0.96	3317

```

[ ]: sns.heatmap(confusion_matrix(y_test, bag_predict), annot=True, fmt='g', _
↳ cmap='Blues')
plt.title("BaggingClassifier")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()

```



```
[ ]: from sklearn.ensemble import RandomForestClassifier

# defining parameter range
param_grid = {
    "n_estimators": [50,100,150,200,250]
}

grid_rfc = GridSearchCV(RandomForestClassifier(), param_grid, refit = True,
    verbose = 3, cv = 10, n_jobs = -1)

# fitting the model for grid search
grid_rfc.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_rfc.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_rfc.best_estimator_)
print(grid_rfc.best_score_)
```

```
Fitting 10 folds for each of 5 candidates, totalling 50 fits
{'n_estimators': 150}
RandomForestClassifier(n_estimators=150)
0.9609705466470111
```

```
[ ]: rfc_model = grid_rfc.best_estimator_
      #rfc_model = rfc.fit(X_train,y_train.values.ravel())
```

```
[ ]: rfc_predict = rfc_model.predict(X_test)
```

```
[ ]: print('The accuracy of RandomForest Classifier is: ' , 100.0 *
      ↪accuracy_score(rfc_predict,y_test))
```


The accuracy of RandomForest Classifier is: 96.8344889960808

```
[ ]: print(classification_report(y_test, rfc_predict))
```


	precision	recall	f1-score	support
0	0.97	0.96	0.96	1444
1	0.97	0.98	0.97	1873
accuracy			0.97	3317
macro avg	0.97	0.97	0.97	3317
weighted avg	0.97	0.97	0.97	3317

```
[ ]: sns.heatmap(confusion_matrix(y_test, rfc_predict), annot=True, fmt='g',
      ↪cmap='Blues')
plt.title("RandomForestClassifier")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
```





```
[ ]: estimators =  
    ↳ [logr_model, knn_model, svc_model, nusvc_model, lsvc_model, xgb_model, ada_model, gbc_model, bag_mo  
  
for estimator in estimators:  
    RocCurveDisplay.from_estimator(estimator, X_test, y_test, ax=plt.gca())
```



```
[ ]: import tensorflow as tf
      #from tensorflow.keras.datasets import imdb
      from keras.layers import Embedding, Dense, LSTM, BatchNormalization
      from keras.losses import BinaryCrossentropy
      from keras.models import Sequential
      from keras.optimizers import Adam
      #from tensorflow.keras.preprocessing.sequence import pad_sequences

      # Model configuration
      additional_metrics = ['accuracy']
      batch_size = 32
      #embedding_output_dims = (X_train.shape[1])
      loss_function = BinaryCrossentropy()
      #max_sequence_length = (X_train.shape[1])
      #num_distinct_words = (X_train.shape[1])
      number_of_epochs = 100
      optimizer = Adam()
      validation_split = 0.20
      verbosity_mode = 1

      # reshape from [samples, features] into [samples, timesteps, features]
```

```

timesteps = 1
X_train_reshape = X_train.values.ravel().reshape(X_train.shape[0],timesteps,
↳X_train.shape[1])
X_test_reshape = X_test.values.ravel().reshape(X_test.shape[0],timesteps,
↳X_test.shape[1])

# Disable eager execution
#tf.compat.v1.disable_eager_execution()

# Load dataset
# (x_train, y_train), (x_test, y_test) = imdb.
↳load_data(num_words=num_distinct_words)
# print(x_train.shape)
# print(x_test.shape)

# Pad all sequences
# padded_inputs = pad_sequences(X_train, maxlen=max_sequence_length, value = 0.
↳0) # 0.0 because it corresponds with <PAD>
# padded_inputs_test = pad_sequences(X_test, maxlen=max_sequence_length, value
↳= 0.0) # 0.0 because it corresponds with <PAD>

# Define the Keras model
def build_model_lstm():
    model = Sequential()
    #model.add(Embedding(num_distinct_words, embedding_output_dims,
↳input_length=max_sequence_length))
    model.add(LSTM(100, input_shape = (timesteps,X_train_reshape.shape[2])))
    model.add(BatchNormalization())
    model.add(Dense(50, activation='relu'))
    model.add(Dense(25, activation='relu'))
    model.add(Dense(10, activation='relu'))
    model.add(Dense(1, activation='sigmoid'))

    # Compile the model
    model.compile(optimizer=optimizer, loss=loss_function,
↳metrics=additional_metrics)
    return model

#from keras.wrappers.scikit_learn import KerasClassifier
lstm_model = build_model_lstm()
# Give a summary
lstm_model.summary()

# Train the model

```

```

history = lstm_model.fit(X_train_reshape, y_train.values.ravel(),
    ↪batch_size=batch_size, epochs=number_of_epochs, verbose=verbosity_mode,
    ↪validation_split=validation_split)

# Test the model after training
#lstm_predict = lstm_model.predict(X_test_reshape)
test_results = lstm_model.evaluate(X_test_reshape, y_test.values.ravel(),
    ↪verbose=False)
print(f'Test results - Loss: {test_results[0]} - Accuracy:
    ↪{100*test_results[1]}%')

```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
lstm_3 (LSTM)	(None, 100)	47600
batch_normalization_3 (Batch Normalization)	(None, 100)	400
lstm_3 (LSTM)	(None, 100)	47600
batch_normalization_3 (Batch Normalization)	(None, 100)	400
dense_12 (Dense)	(None, 50)	5050
dense_13 (Dense)	(None, 25)	1275
dense_14 (Dense)	(None, 10)	260
dense_15 (Dense)	(None, 1)	11

Total params: 54,596  
 Trainable params: 54,396  
 Non-trainable params: 200

Epoch 1/100  
 194/194 [=====] - 5s 8ms/step - loss: 0.2603 - accuracy: 0.8955 - val\_loss: 0.4548 - val\_accuracy: 0.8579  
 Epoch 2/100  
 194/194 [=====] - 1s 5ms/step - loss: 0.1773 - accuracy: 0.9299 - val\_loss: 0.2267 - val\_accuracy: 0.9451  
 Epoch 3/100

194/194 [=====] - 1s 5ms/step - loss: 0.1618 -  
accuracy: 0.9341 - val\_loss: 0.1524 - val\_accuracy: 0.9516  
Epoch 4/100  
194/194 [=====] - 1s 5ms/step - loss: 0.1467 -  
accuracy: 0.9376 - val\_loss: 0.1359 - val\_accuracy: 0.9419  
Epoch 5/100  
194/194 [=====] - 1s 5ms/step - loss: 0.1452 -  
accuracy: 0.9402 - val\_loss: 0.1270 - val\_accuracy: 0.9451  
Epoch 6/100  
194/194 [=====] - 1s 4ms/step - loss: 0.1296 -  
accuracy: 0.9465 - val\_loss: 0.1407 - val\_accuracy: 0.9354  
Epoch 7/100  
194/194 [=====] - 1s 4ms/step - loss: 0.1331 -  
accuracy: 0.9435 - val\_loss: 0.1259 - val\_accuracy: 0.9399  
Epoch 8/100  
194/194 [=====] - 1s 5ms/step - loss: 0.1217 -  
accuracy: 0.9504 - val\_loss: 0.1272 - val\_accuracy: 0.9444  
Epoch 9/100  
194/194 [=====] - 1s 4ms/step - loss: 0.1158 -  
accuracy: 0.9501 - val\_loss: 0.1323 - val\_accuracy: 0.9386  
Epoch 10/100  
194/194 [=====] - 1s 4ms/step - loss: 0.1154 -  
accuracy: 0.9494 - val\_loss: 0.1283 - val\_accuracy: 0.9503  
Epoch 11/100  
194/194 [=====] - 1s 5ms/step - loss: 0.1107 -  
accuracy: 0.9527 - val\_loss: 0.1281 - val\_accuracy: 0.9457  
Epoch 12/100  
194/194 [=====] - 1s 4ms/step - loss: 0.1003 -  
accuracy: 0.9561 - val\_loss: 0.1134 - val\_accuracy: 0.9490  
Epoch 13/100  
194/194 [=====] - 1s 6ms/step - loss: 0.0935 -  
accuracy: 0.9603 - val\_loss: 0.1340 - val\_accuracy: 0.9419  
Epoch 14/100  
194/194 [=====] - 1s 8ms/step - loss: 0.1042 -  
accuracy: 0.9551 - val\_loss: 0.1447 - val\_accuracy: 0.9438  
Epoch 15/100  
194/194 [=====] - 1s 7ms/step - loss: 0.0953 -  
accuracy: 0.9588 - val\_loss: 0.1218 - val\_accuracy: 0.9464  
Epoch 16/100  
194/194 [=====] - 1s 6ms/step - loss: 0.0994 -  
accuracy: 0.9588 - val\_loss: 0.1106 - val\_accuracy: 0.9490  
Epoch 17/100  
194/194 [=====] - 1s 6ms/step - loss: 0.0910 -  
accuracy: 0.9624 - val\_loss: 0.1231 - val\_accuracy: 0.9496  
Epoch 18/100  
194/194 [=====] - 1s 6ms/step - loss: 0.0880 -  
accuracy: 0.9620 - val\_loss: 0.1262 - val\_accuracy: 0.9528  
Epoch 19/100

194/194 [=====] - 1s 6ms/step - loss: 0.0832 -  
accuracy: 0.9640 - val\_loss: 0.1258 - val\_accuracy: 0.9509  
Epoch 20/100  
194/194 [=====] - 1s 6ms/step - loss: 0.0864 -  
accuracy: 0.9624 - val\_loss: 0.1364 - val\_accuracy: 0.9477  
Epoch 21/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0832 -  
accuracy: 0.9653 - val\_loss: 0.1220 - val\_accuracy: 0.9490  
Epoch 22/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0866 -  
accuracy: 0.9633 - val\_loss: 0.1320 - val\_accuracy: 0.9483  
Epoch 23/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0752 -  
accuracy: 0.9690 - val\_loss: 0.1219 - val\_accuracy: 0.9470  
Epoch 24/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0744 -  
accuracy: 0.9680 - val\_loss: 0.1222 - val\_accuracy: 0.9496  
Epoch 25/100  
194/194 [=====] - 1s 6ms/step - loss: 0.0748 -  
accuracy: 0.9701 - val\_loss: 0.1297 - val\_accuracy: 0.9516  
Epoch 26/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0741 -  
accuracy: 0.9683 - val\_loss: 0.1276 - val\_accuracy: 0.9535  
Epoch 27/100  
194/194 [=====] - 1s 6ms/step - loss: 0.0713 -  
accuracy: 0.9703 - val\_loss: 0.1350 - val\_accuracy: 0.9451  
Epoch 28/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0704 -  
accuracy: 0.9722 - val\_loss: 0.1302 - val\_accuracy: 0.9483  
Epoch 29/100  
194/194 [=====] - 1s 6ms/step - loss: 0.0742 -  
accuracy: 0.9696 - val\_loss: 0.1379 - val\_accuracy: 0.9425  
Epoch 30/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0740 -  
accuracy: 0.9679 - val\_loss: 0.1553 - val\_accuracy: 0.9354  
Epoch 31/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0651 -  
accuracy: 0.9721 - val\_loss: 0.1349 - val\_accuracy: 0.9477  
Epoch 32/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0708 -  
accuracy: 0.9691 - val\_loss: 0.1258 - val\_accuracy: 0.9528  
Epoch 33/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0663 -  
accuracy: 0.9712 - val\_loss: 0.1260 - val\_accuracy: 0.9561  
Epoch 34/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0741 -  
accuracy: 0.9682 - val\_loss: 0.1424 - val\_accuracy: 0.9496  
Epoch 35/100

194/194 [=====] - 1s 5ms/step - loss: 0.0645 -  
accuracy: 0.9742 - val\_loss: 0.1378 - val\_accuracy: 0.9516  
Epoch 36/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0690 -  
accuracy: 0.9709 - val\_loss: 0.1287 - val\_accuracy: 0.9503  
Epoch 37/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0603 -  
accuracy: 0.9758 - val\_loss: 0.1341 - val\_accuracy: 0.9516  
Epoch 38/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0611 -  
accuracy: 0.9769 - val\_loss: 0.1271 - val\_accuracy: 0.9554  
Epoch 39/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0636 -  
accuracy: 0.9704 - val\_loss: 0.1490 - val\_accuracy: 0.9541  
Epoch 40/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0621 -  
accuracy: 0.9740 - val\_loss: 0.1272 - val\_accuracy: 0.9561  
Epoch 41/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0624 -  
accuracy: 0.9725 - val\_loss: 0.1163 - val\_accuracy: 0.9541  
Epoch 42/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0628 -  
accuracy: 0.9740 - val\_loss: 0.1324 - val\_accuracy: 0.9554  
Epoch 43/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0600 -  
accuracy: 0.9767 - val\_loss: 0.1351 - val\_accuracy: 0.9528  
Epoch 44/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0584 -  
accuracy: 0.9751 - val\_loss: 0.1457 - val\_accuracy: 0.9528  
Epoch 45/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0557 -  
accuracy: 0.9761 - val\_loss: 0.1404 - val\_accuracy: 0.9522  
Epoch 46/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0577 -  
accuracy: 0.9746 - val\_loss: 0.1396 - val\_accuracy: 0.9561  
Epoch 47/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0567 -  
accuracy: 0.9742 - val\_loss: 0.1508 - val\_accuracy: 0.9451  
Epoch 48/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0584 -  
accuracy: 0.9751 - val\_loss: 0.1470 - val\_accuracy: 0.9470  
Epoch 49/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0560 -  
accuracy: 0.9772 - val\_loss: 0.1479 - val\_accuracy: 0.9457  
Epoch 50/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0535 -  
accuracy: 0.9764 - val\_loss: 0.1544 - val\_accuracy: 0.9470  
Epoch 51/100

194/194 [=====] - 1s 4ms/step - loss: 0.0578 -  
accuracy: 0.9738 - val\_loss: 0.1512 - val\_accuracy: 0.9528  
Epoch 52/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0558 -  
accuracy: 0.9751 - val\_loss: 0.1437 - val\_accuracy: 0.9477  
Epoch 53/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0527 -  
accuracy: 0.9785 - val\_loss: 0.1440 - val\_accuracy: 0.9522  
Epoch 54/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0508 -  
accuracy: 0.9805 - val\_loss: 0.1612 - val\_accuracy: 0.9496  
Epoch 55/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0536 -  
accuracy: 0.9785 - val\_loss: 0.1467 - val\_accuracy: 0.9574  
Epoch 56/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0565 -  
accuracy: 0.9759 - val\_loss: 0.1569 - val\_accuracy: 0.9541  
Epoch 57/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0529 -  
accuracy: 0.9771 - val\_loss: 0.1508 - val\_accuracy: 0.9432  
Epoch 58/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0493 -  
accuracy: 0.9793 - val\_loss: 0.1402 - val\_accuracy: 0.9541  
Epoch 59/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0470 -  
accuracy: 0.9790 - val\_loss: 0.1453 - val\_accuracy: 0.9561  
Epoch 60/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0464 -  
accuracy: 0.9792 - val\_loss: 0.1463 - val\_accuracy: 0.9535  
Epoch 61/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0533 -  
accuracy: 0.9753 - val\_loss: 0.1587 - val\_accuracy: 0.9574  
Epoch 62/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0557 -  
accuracy: 0.9771 - val\_loss: 0.1496 - val\_accuracy: 0.9483  
Epoch 63/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0492 -  
accuracy: 0.9787 - val\_loss: 0.1446 - val\_accuracy: 0.9522  
Epoch 64/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0515 -  
accuracy: 0.9775 - val\_loss: 0.1737 - val\_accuracy: 0.9528  
Epoch 65/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0510 -  
accuracy: 0.9759 - val\_loss: 0.1396 - val\_accuracy: 0.9509  
Epoch 66/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0472 -  
accuracy: 0.9788 - val\_loss: 0.1553 - val\_accuracy: 0.9528  
Epoch 67/100



194/194 [=====] - 1s 4ms/step - loss: 0.0491 -  
accuracy: 0.9788 - val\_loss: 0.1656 - val\_accuracy: 0.9522  
Epoch 68/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0474 -  
accuracy: 0.9801 - val\_loss: 0.1542 - val\_accuracy: 0.9528  
Epoch 69/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0483 -  
accuracy: 0.9803 - val\_loss: 0.1530 - val\_accuracy: 0.9496  
Epoch 70/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0456 -  
accuracy: 0.9801 - val\_loss: 0.1645 - val\_accuracy: 0.9522  
Epoch 71/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0458 -  
accuracy: 0.9795 - val\_loss: 0.1596 - val\_accuracy: 0.9490  
Epoch 72/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0470 -  
accuracy: 0.9801 - val\_loss: 0.1647 - val\_accuracy: 0.9496  
Epoch 73/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0457 -  
accuracy: 0.9771 - val\_loss: 0.1519 - val\_accuracy: 0.9503  
Epoch 74/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0461 -  
accuracy: 0.9800 - val\_loss: 0.1587 - val\_accuracy: 0.9548  
Epoch 75/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0448 -  
accuracy: 0.9798 - val\_loss: 0.1600 - val\_accuracy: 0.9483  
Epoch 76/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0480 -  
accuracy: 0.9764 - val\_loss: 0.1725 - val\_accuracy: 0.9522  
Epoch 77/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0475 -  
accuracy: 0.9790 - val\_loss: 0.1566 - val\_accuracy: 0.9541  
Epoch 78/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0435 -  
accuracy: 0.9805 - val\_loss: 0.1616 - val\_accuracy: 0.9535  
Epoch 79/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0447 -  
accuracy: 0.9811 - val\_loss: 0.1799 - val\_accuracy: 0.9535  
Epoch 80/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0458 -  
accuracy: 0.9800 - val\_loss: 0.1580 - val\_accuracy: 0.9541  
Epoch 81/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0499 -  
accuracy: 0.9772 - val\_loss: 0.1609 - val\_accuracy: 0.9490  
Epoch 82/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0443 -  
accuracy: 0.9800 - val\_loss: 0.1803 - val\_accuracy: 0.9509  
Epoch 83/100

194/194 [=====] - 1s 4ms/step - loss: 0.0449 -  
accuracy: 0.9801 - val\_loss: 0.1640 - val\_accuracy: 0.9464  
Epoch 84/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0384 -  
accuracy: 0.9821 - val\_loss: 0.1629 - val\_accuracy: 0.9516  
Epoch 85/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0379 -  
accuracy: 0.9824 - val\_loss: 0.1629 - val\_accuracy: 0.9528  
Epoch 86/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0475 -  
accuracy: 0.9779 - val\_loss: 0.1625 - val\_accuracy: 0.9464  
Epoch 87/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0465 -  
accuracy: 0.9785 - val\_loss: 0.1730 - val\_accuracy: 0.9516  
Epoch 88/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0485 -  
accuracy: 0.9795 - val\_loss: 0.1710 - val\_accuracy: 0.9535  
Epoch 89/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0454 -  
accuracy: 0.9793 - val\_loss: 0.1636 - val\_accuracy: 0.9490  
Epoch 90/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0418 -  
accuracy: 0.9811 - val\_loss: 0.1861 - val\_accuracy: 0.9522  
Epoch 91/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0431 -  
accuracy: 0.9801 - val\_loss: 0.1836 - val\_accuracy: 0.9548  
Epoch 92/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0435 -  
accuracy: 0.9788 - val\_loss: 0.1639 - val\_accuracy: 0.9567  
Epoch 93/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0417 -  
accuracy: 0.9811 - val\_loss: 0.1694 - val\_accuracy: 0.9503  
Epoch 94/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0410 -  
accuracy: 0.9817 - val\_loss: 0.1802 - val\_accuracy: 0.9516  
Epoch 95/100  
194/194 [=====] - 1s 5ms/step - loss: 0.0402 -  
accuracy: 0.9805 - val\_loss: 0.1924 - val\_accuracy: 0.9567  
Epoch 96/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0443 -  
accuracy: 0.9792 - val\_loss: 0.1602 - val\_accuracy: 0.9593  
Epoch 97/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0468 -  
accuracy: 0.9788 - val\_loss: 0.1883 - val\_accuracy: 0.9522  
Epoch 98/100  
194/194 [=====] - 1s 4ms/step - loss: 0.0421 -  
accuracy: 0.9808 - val\_loss: 0.1614 - val\_accuracy: 0.9522  
Epoch 99/100

```

194/194 [=====] - 1s 4ms/step - loss: 0.0389 -
accuracy: 0.9821 - val_loss: 0.1944 - val_accuracy: 0.9535
Epoch 100/100
194/194 [=====] - 1s 4ms/step - loss: 0.0411 -
accuracy: 0.9817 - val_loss: 0.1678 - val_accuracy: 0.9522
Test results - Loss: 0.15201237797737122 - Accuracy: 95.47784328460693%

```

```

[ ]: lstm_predict_proba = lstm_model.predict(X_test_reshape, batch_size=32)
lstm_predict_class = (lstm_predict_proba > 0.5).astype("int32")
print(classification_report(y_test, lstm_predict_class))

```

```

104/104 [=====] - 1s 2ms/step


```

	precision	recall	f1-score	support
0	0.95	0.94	0.95	1444
1	0.96	0.96	0.96	1873
accuracy			0.95	3317
macro avg	0.95	0.95	0.95	3317
weighted avg	0.95	0.95	0.95	3317

```

[ ]: sns.heatmap(confusion_matrix(y_test, lstm_predict_class), annot=True, fmt='g',
cmap='Blues')
plt.title("LSTM")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()

```



```
[ ]: RocCurveDisplay.from_predictions(y_test,lstm_predict_class)
plt.show()
```

