

## chi\_sq\_30 9010 split .05 threshold

January 2, 2023

```
[ ]: # Importing the packages
import sys
import numpy as np
np.set_printoptions(threshold=sys.maxsize)
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
import sklearn
import random
from sklearn.metrics import ↵
    ↵confusion_matrix, accuracy_score, classification_report, RocCurveDisplay, ConfusionMatrixDisplay

[ ]: pd.set_option('display.max_rows', None)
pd.set_option('display.max_columns', None)
pd.set_option('display.width', None)
pd.set_option('display.max_colwidth', None)

[ ]: # Importing the dataset
df = pd.read_csv('dataset_30.csv')
df.drop(['index'], axis=1, inplace=True)
#df.head()

[ ]: # if your dataset contains missing value, check which column has missing values
#df.isnull().sum()

[ ]: #df.dropna(inplace=True)

[ ]: from sklearn import preprocessing

col = df.columns[:]

lab_en= preprocessing.LabelEncoder()

for c in col:
    df[c]= lab_en.fit_transform(df[c])

#df.head(50)
```

```
[ ]: a=len(df[df.Result==0])
      b=len(df[df.Result==1])
```

```
[ ]: print("Count of Legitimate Websites = ", a)
      print("Count of Phishy Websites = ", b)
```

Count of Legitimate Websites = 4898  
Count of Phishy Websites = 6157

```
[ ]: X = df.drop(['Result'], axis=1, inplace=False)
      #X.head()
      #same work
      ##inplace true modifies the og data & does not return anything
      ##inplace false does not modify og data but returns something which we store in
      ↪ a var
      # X= df.drop(columns='Result')
      # X.head()
```

```
[ ]: #df.head()
```

```
[ ]: y = df['Result']
      y = pd.DataFrame(y)
      y.head()
```

```
[ ]:      Result
      0      0
      1      0
      2      0
      3      0
      4      1
```

```
[ ]: # separate dataset into train and test
      from cProfile import label
      from sklearn.model_selection import train_test_split
      X_train, X_test, y_train, y_test = train_test_split(
          X,
          y,
          test_size=0.1,
          random_state=10)

      X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
[ ]: ((9949, 30), (1106, 30), (9949, 1), (1106, 1))
```

```
[ ]: #perform chi square test
      from sklearn.feature_selection import chi2
      f_p_values = chi2(X_train,y_train)
```

```
[ ]: f_p_values
```

```
[ ]: (array([2.95454443e+01, 5.69285338e+01, 6.48054923e+00, 5.19005883e+00,
          2.24367695e+00, 1.05380952e+03, 5.83898807e+02, 3.39307396e+03,
          3.43323680e+02, 7.10402361e-03, 2.19169758e+00, 3.28375268e+00,
          2.57362852e+02, 2.67851850e+03, 4.16627163e+02, 6.85589281e+02,
          9.97960542e-01, 6.39645595e+00, 2.83434016e+00, 2.37822783e+00,
          9.20405451e-02, 6.23476770e-03, 1.16281643e-03, 6.87102345e+01,
          1.79259515e+01, 6.19393719e+02, 8.66397569e+01, 2.26984316e+01,
          2.81600035e+00, 8.30737109e+00]),
      array([5.46208850e-08, 4.51940623e-014, 1.09061279e-002, 2.27164473e-002,
          1.34161615e-001, 3.61673986e-231, 5.32127479e-129, 0.00000000e+000,
          1.20515364e-076, 9.32829544e-001, 1.38756314e-001, 6.99687740e-002,
          6.44707289e-058, 0.00000000e+000, 1.32309547e-092, 4.06870428e-151,
          3.17804501e-001, 1.14348410e-002, 9.22686941e-002, 1.23037052e-001,
          7.61598878e-001, 9.37064005e-001, 9.72797333e-001, 1.14047515e-016,
          2.29668002e-005, 1.01302745e-136, 1.30203200e-020, 1.89522496e-006,
          9.33286991e-002, 3.94845012e-003]))
```

```
[ ]: #The less the p_values the more important that feature is
p_values = pd.Series(f_p_values[1])
p_values.index = X_train.columns
p_values
```

```
[ ]: having_IPhaving_IP_Address      5.462088e-08
URLURL_Length                      4.519406e-14
Shortining_Service                 1.090613e-02
having_At_Symbol                   2.271645e-02
double_slash_redirecting           1.341616e-01
Prefix_Suffix                     3.616740e-231
having_Sub_Domain                  5.321275e-129
SSLfinal_State                     0.000000e+00
Domain_registration_length         1.205154e-76
Favicon                           9.328295e-01
port                               1.387563e-01
HTTPS_token                        6.996877e-02
Request_URL                        6.447073e-58
URL_of_Anchor                      0.000000e+00
Links_in_tags                      1.323095e-92
SFH                                4.068704e-151
Submitting_to_email                 3.178045e-01
Abnormal_URL                       1.143484e-02
Redirect                           9.226869e-02
on_mouseover                       1.230371e-01
RightClick                         7.615989e-01
popUpWidnow                        9.370640e-01
Iframe                             9.727973e-01
```

```

age_of_domain          1.140475e-16
DNSRecord              2.296680e-05
web_traffic            1.013027e-136
Page_Rank              1.302032e-20
Google_Index          1.895225e-06
Links_pointing_to_page 9.332870e-02
Statistical_report     3.948450e-03
dtype: float64

```

```

[ ]: #sort p_values to check which feature has the lowest values
p_values = p_values.sort_values(ascending = False)
p_values

```

```

[ ]: Iframe          9.727973e-01
popUpWidnow         9.370640e-01
Favicon             9.328295e-01
RightClick          7.615989e-01
Submitting_to_email 3.178045e-01
port                1.387563e-01
double_slash_redirecting 1.341616e-01
on_mouseover        1.230371e-01
Links_pointing_to_page 9.332870e-02
Redirect            9.226869e-02
HTTPS_token         6.996877e-02
having_At_Symbol    2.271645e-02
Abnormal_URL        1.143484e-02
Shortining_Service  1.090613e-02
Statistical_report  3.948450e-03
DNSRecord           2.296680e-05
Google_Index        1.895225e-06
having_IPhaving_IP_Address 5.462088e-08
URLURL_Length       4.519406e-14
age_of_domain       1.140475e-16
Page_Rank           1.302032e-20
Request_URL         6.447073e-58
Domain_registration_length 1.205154e-76
Links_in_tags       1.323095e-92
having_Sub_Domain   5.321275e-129
web_traffic         1.013027e-136
SFH                 4.068704e-151
Prefix_Suffix       3.616740e-231
URL_of_Anchor       0.000000e+00
SSLfinal_State      0.000000e+00
dtype: float64

```

```

[ ]: def DropFeature (p_values, threshold):
      drop_feature = set()

```

```

    for index, values in p_values.items():
        if values > threshold or np.isnan(values):
            drop_feature.add(index)
    return drop_feature

```

```

[ ]: drop_feature = DropFeature(p_values, .05)
    len(set(drop_feature))

```

```

[ ]: 11

```

```

[ ]: drop_feature

```

```

[ ]: {'Favicon',
      'HTTPS_token',
      'Iframe',
      'Links_pointing_to_page',
      'Redirect',
      'RightClick',
      'Submitting_to_email',
      'double_slash_redirecting',
      'on_mouseover',
      'popUpWidnow',
      'port'}

```

```

[ ]: X_train.drop(drop_feature, axis=1, inplace=True)
    X_test.drop(drop_feature, axis=1, inplace=True)

```

```

[ ]: len(df.columns)

```

```

[ ]: 31

```

```

[ ]: print("Training set has {} samples.".format(X_train.shape[0]))
    print("Testing set has {} samples.".format(X_test.shape[0]))

```

Training set has 9949 samples.

Testing set has 1106 samples.

```

[ ]: from sklearn.model_selection import GridSearchCV
    from sklearn.linear_model import LogisticRegression

    # defining parameter range
    param_grid = {'penalty' : ['l2'],
                  'C' : [0.1, 1, 10, 20, 30],
                  'solver' : ['lbfgs', 'newton-cg', 'liblinear', 'sag', 'saga'],
                  'max_iter' : [2500, 5000]}

    grid_logr = GridSearchCV(LogisticRegression(), param_grid, refit = True, cv = 10,
                             verbose = 3, n_jobs = -1)

```

```

# fitting the model for grid search
grid_logr.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_logr.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_logr.best_estimator_)
print(grid_logr.best_score_)

```

```

Fitting 10 folds for each of 50 candidates, totalling 500 fits
{'C': 10, 'max_iter': 2500, 'penalty': 'l2', 'solver': 'lbfgs'}
LogisticRegression(C=10, max_iter=2500)
0.9239128236757226

```

```

[ ]: logr_model = grid_logr.best_estimator_

# Performing training
#logr_model = logr.fit(X_train, y_train.values.ravel())

```

```

[ ]: logr_predict = logr_model.predict(X_test)

```

```

[ ]: # from sklearn.metrics import confusion_matrix, accuracy_score
# cm = confusion_matrix(y_test, dct_pred)
# ac = accuracy_score(y_test, dct_pred)

```

```

[ ]: print ("Accuracy of logr classifier : ", accuracy_score(y_test,
↳logr_predict)*100)

```

```

Accuracy of logr classifier : 92.6763110307414

```

```

[ ]: print(classification_report(y_test, logr_predict))

```


	precision	recall	f1-score	support
0	0.93	0.90	0.91	472
1	0.93	0.95	0.94	634
accuracy			0.93	1106
macro avg	0.93	0.92	0.92	1106
weighted avg	0.93	0.93	0.93	1106

```

[ ]: sns.heatmap(confusion_matrix(y_test, logr_predict), annot=True, fmt='g',
↳ cmap='Blues')
plt.title("LogisticRegression")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')

```

```
plt.show()
```



```
[ ]: # from sklearn.neighbors import KNeighborsClassifier

# #training_accuracy=[]
# test_accuracy=[]

# neighbors=range(1,10)
# ##values.ravel() converts vector y to flattened array
# for i in neighbors:
#     knn=KNeighborsClassifier(n_neighbors=i)
#     knn_model = knn.fit(X_train,y_train.values.ravel())
#     #training_accuracy.append(knn.score(X_train,y_train.values.ravel()))
#     test_accuracy.append(knn_model.score(X_test,y_test.values.ravel()))
```

```
[ ]: # plt.plot(neighbors,test_accuracy,label="test accuracy")
# plt.ylabel("Accuracy")
# plt.xlabel("number of neighbors")
# plt.legend()
```

```
# plt.show()
```

```
[ ]: from sklearn.neighbors import KNeighborsClassifier

# defining parameter range
param_grid = {'n_neighbors': [1,2,3,4,5,6,7,8,9,10]}

grid_knn = GridSearchCV(KNeighborsClassifier(), param_grid, refit = True, cv = 10, verbose = 3, n_jobs = -1)

# fitting the model for grid search
grid_knn.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_knn.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_knn.best_estimator_)
print(grid_knn.best_score_)
```

Fitting 10 folds for each of 10 candidates, totalling 100 fits

```
{'n_neighbors': 1}
```

```
KNeighborsClassifier(n_neighbors=1)
```

```
0.9551719361394498
```

```
[ ]: knn_model = grid_knn.best_estimator_
#knn_model = knn.fit(X_train,y_train.values.ravel())
```

```
[ ]: #print ("Accuracy of knn classifier: ", max(test_accuracy)*100)
knn_predict = knn_model.predict(X_test)
```

```
[ ]: print('The accuracy of knn Classifier is: ', 100.0 * accuracy_score(y_test, knn_predict))
```


The accuracy of knn Classifier is: 96.29294755877035

```
[ ]: print(classification_report(y_test, knn_predict))
```

	precision	recall	f1-score	support
0	0.97	0.94	0.96	472
1	0.96	0.98	0.97	634
accuracy			0.96	1106
macro avg	0.96	0.96	0.96	1106
weighted avg	0.96	0.96	0.96	1106



```
[ ]: sns.heatmap(confusion_matrix(y_test, knn_predict), annot=True, fmt='g',
    cmap='Blues')
plt.title("KNeighborsClassifier")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
```



```
[ ]: ## here is the change
# knn_y_pred_proba = knn.predict_proba(X_test)
# knn_y_pred_proba_positive = knn_y_pred_proba[:, 1]

# RocCurveDisplay.from_predictions(y_test, knn_y_pred_proba_positive)

# fig, ax = plt.subplots()
# RocCurveDisplay.from_estimator(
#     logreg, X_test, y_test, ax = ax)

# logreg_y_decision = logreg.decision_function(X_test)
```

```
# metrics.RocCurveDisplay.
↪from_predictions(y_test, logreg_y_decision, ax=ax, name="logreg predictions")
```

```
[ ]: from sklearn.svm import SVC

# defining parameter range
param_grid = {'C': [0.1, 1, 10],
              'gamma': [1, 0.1, 0.01],
              'kernel': ['linear', 'poly', 'rbf', 'sigmoid']}

grid_svc = GridSearchCV(SVC(), param_grid, refit = True, cv = 10, verbose = 3, ↪
↪n_jobs = -1)

# fitting the model for grid search
grid_svc.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_svc.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_svc.best_estimator_)
print(grid_svc.best_score_)
```

Fitting 10 folds for each of 36 candidates, totalling 360 fits  
{'C': 1, 'gamma': 1, 'kernel': 'rbf'}  
SVC(C=1, gamma=1)  
0.9617048016743677

```
[ ]: svc_model = grid_svc.best_estimator_
#svc_model = svc.fit(X_train, y_train.values.ravel())
```

```
[ ]: svc_predict = svc_model.predict(X_test)
```


```
[ ]: print('The accuracy of svc Classifier is: ', 100.0 * accuracy_score(y_test, ↪
↪svc_predict))
```

The accuracy of svc Classifier is: 96.47377938517178

```
[ ]: print(classification_report(y_test, svc_predict))
```

	precision	recall	f1-score	support
0	0.97	0.95	0.96	472
1	0.96	0.98	0.97	634
accuracy			0.96	1106
macro avg	0.97	0.96	0.96	1106
weighted avg	0.96	0.96	0.96	1106

```
[ ]: sns.heatmap(confusion_matrix(y_test, svc_predict), annot=True, fmt='g',  
    cmap='Blues')  
plt.title("SVC")  
plt.xlabel('Predicted Class')  
plt.ylabel('Original Class')  
plt.show()
```



```
[ ]: from sklearn.svm import NuSVC  
  
# defining parameter range  
param_grid = {'nu': [0.1, 0.5],  
              'gamma': [1, 0.1, 0.01],  
              'kernel': ['linear', 'poly', 'rbf', 'sigmoid']}  
  
grid_nusvc = GridSearchCV(NuSVC(), param_grid, refit = True, verbose = 3, cv =  
    10, n_jobs = -1)  
  
# fitting the model for grid search  
grid_nusvc.fit(X_train, y_train.values.ravel())
```

```
# print best parameter after tuning
print(grid_nusvc.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_nusvc.best_estimator_)
print(grid_nusvc.best_score_)
```

```
Fitting 10 folds for each of 24 candidates, totalling 240 fits
{'gamma': 1, 'kernel': 'rbf', 'nu': 0.1}
NuSVC(gamma=1, nu=0.1)
0.9616039958343021
```

```
[ ]: nusvc_model = grid_nusvc.best_estimator_
#nusvc_model = nusvc.fit(X_train, y_train.values.ravel())
```

```
[ ]: nusvc_predict = nusvc_model.predict(X_test)
```


```
[ ]: print('The accuracy of nusvc Classifier is: ', 100.0 * accuracy_score(y_test,
↪nusvc_predict))
```

```
The accuracy of nusvc Classifier is: 96.74502712477397
```

```
[ ]: print(classification_report(y_test, nusvc_predict))
```

	precision	recall	f1-score	support
0	0.97	0.95	0.96	472
1	0.96	0.98	0.97	634
accuracy			0.97	1106
macro avg	0.97	0.97	0.97	1106
weighted avg	0.97	0.97	0.97	1106

```
[ ]: sns.heatmap(confusion_matrix(y_test, nusvc_predict), annot=True, fmt='g',
↪cmap='Blues')
plt.title("NuSVC")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
```



```
[ ]: from sklearn.svm import LinearSVC

# defining parameter range
param_grid = {'C': [0.1, 1, 10, 20, 30],
              'penalty': ['l1', 'l2'],
              'loss': ['squared_hinge'],
              'dual': [False],
              'tol': [.1, .01, .001]}

grid_lsvc = GridSearchCV(LinearSVC(), param_grid, refit = True, verbose = 3, cv=
↳ 10, n_jobs = -1)

# fitting the model for grid search
grid_lsvc.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_lsvc.best_params_)

# print how our model looks after hyper-parameter tuning
```

```
print(grid_lsvc.best_estimator_)
print(grid_lsvc.best_score_)
```

Fitting 10 folds for each of 30 candidates, totalling 300 fits  
{'C': 1, 'dual': False, 'loss': 'squared\_hinge', 'penalty': 'l2', 'tol': 0.01}  
LinearSVC(C=1, dual=False, tol=0.01)  
0.924616644591165

```
[ ]: lsvc_model = grid_lsvc.best_estimator_
      #lsvc_model = lsvc.fit(X_train, y_train.values.ravel())
```

```
[ ]: lsvc_predict = lsvc_model.predict(X_test)
```

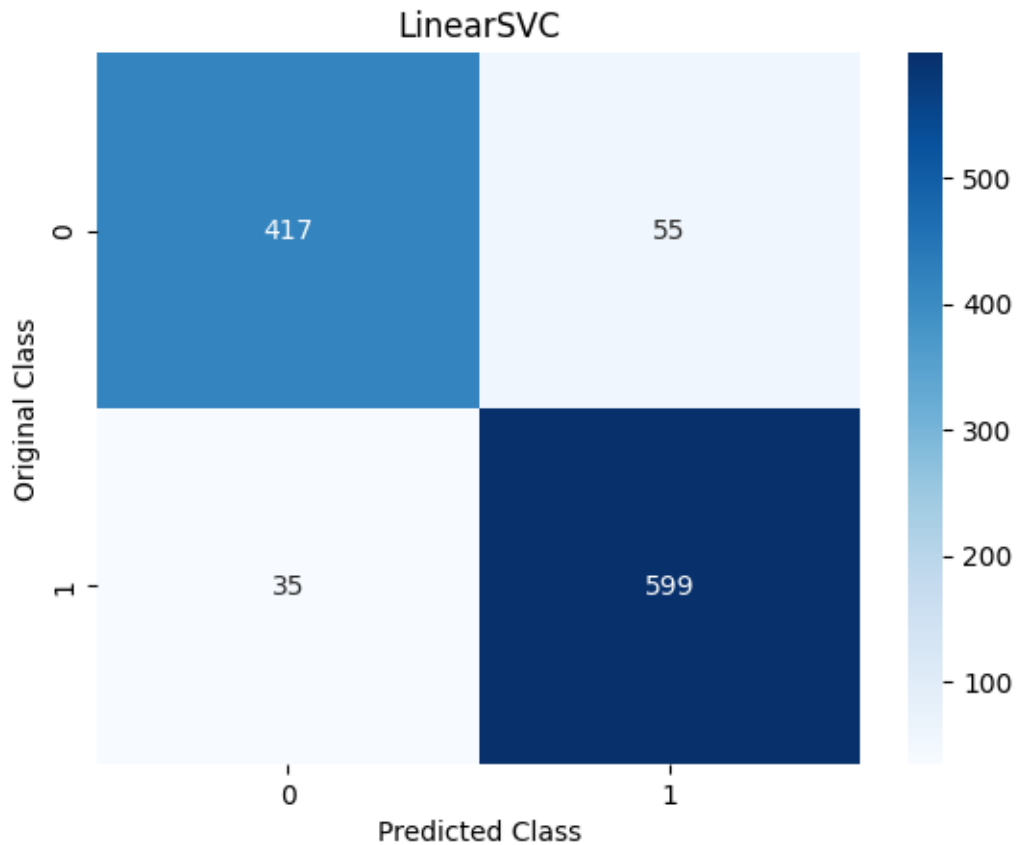
```
[ ]: print('The accuracy of lsvc Classifier is: ', 100.0 * accuracy_score(y_test, lsvc_predict))
```

The accuracy of lsvc Classifier is: 91.86256781193491

```
[ ]: print(classification_report(y_test, lsvc_predict))
```

	precision	recall	f1-score	support
0	0.92	0.88	0.90	472
1	0.92	0.94	0.93	634
accuracy			0.92	1106
macro avg	0.92	0.91	0.92	1106
weighted avg	0.92	0.92	0.92	1106

```
[ ]: sns.heatmap(confusion_matrix(y_test, lsvc_predict), annot=True, fmt='g', cmap='Blues')
      plt.title("LinearSVC")
      plt.xlabel('Predicted Class')
      plt.ylabel('Original Class')
      plt.show()
```



```
[ ]: from sklearn.ensemble import AdaBoostClassifier

# defining parameter range
param_grid = {'n_estimators': [40,50,100,200,300]}

grid_ada = GridSearchCV(AdaBoostClassifier(), param_grid, refit = True, verbose=
    ↪ 3, cv = 10, n_jobs = -1)

# fitting the model for grid search
grid_ada.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_ada.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_ada.best_estimator_)
print(grid_ada.best_score_)
```

Fitting 10 folds for each of 5 candidates, totalling 50 fits  
 {'n\_estimators': 100}

```
AdaBoostClassifier(n_estimators=100)
0.9330590578647767
```

```
[ ]: ada_model = grid_ada.best_estimator_  
      #ada_model = ada.fit(X_train,y_train.values.ravel())
```

```
[ ]: ada_predict = ada_model.predict(X_test)
```

```
[ ]: print('The accuracy of Ada Boost Classifier is: ', 100.0 *  
      ↪accuracy_score(ada_predict,y_test))
```

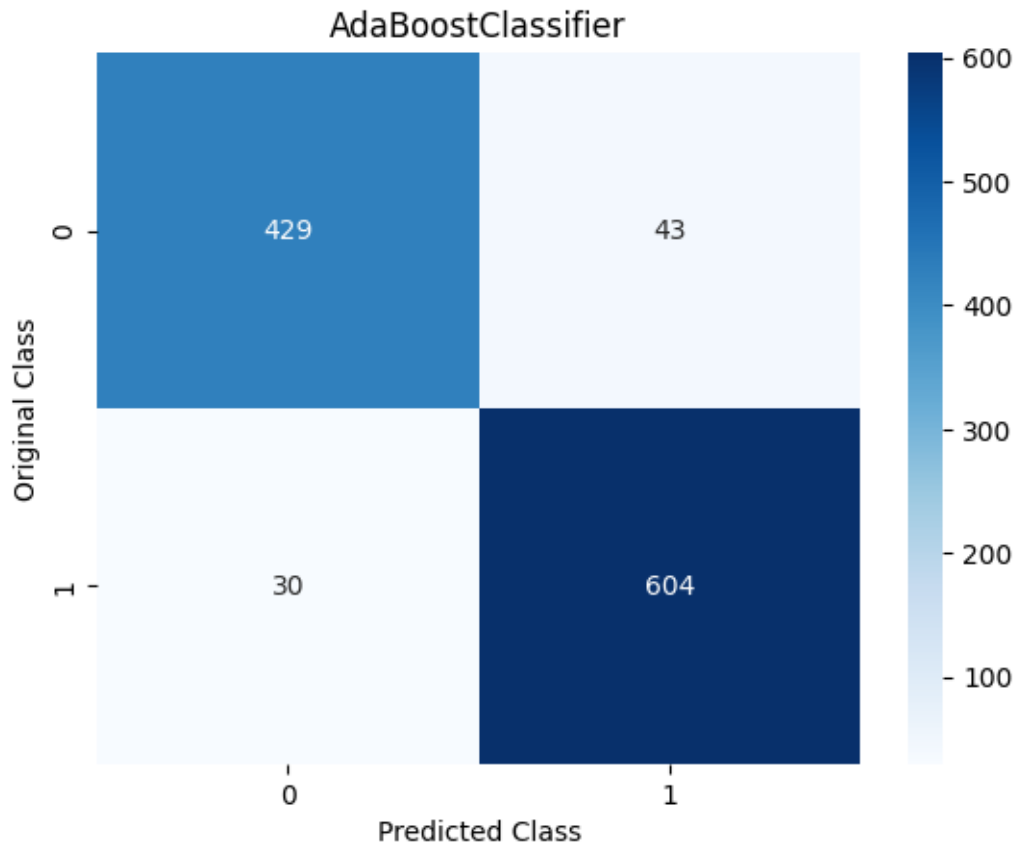
The accuracy of Ada Boost Classifier is: 93.3996383363472

```
[ ]: print(classification_report(y_test, ada_predict))
```

	precision	recall	f1-score	support
0	0.93	0.91	0.92	472
1	0.93	0.95	0.94	634
accuracy			0.93	1106
macro avg	0.93	0.93	0.93	1106
weighted avg	0.93	0.93	0.93	1106

```
[ ]: sns.heatmap(confusion_matrix(y_test, ada_predict), annot=True, fmt='g',  
      ↪cmap='Blues')  
plt.title("AdaBoostClassifier")  
plt.xlabel('Predicted Class')  
plt.ylabel('Original Class')  
plt.show()
```





```
[ ]: from xgboost import XGBClassifier

# defining parameter range
param_grid = {
    "gamma": [.01, .1, .5],
    "n_estimators": [50,100,150,200,250]
}

grid_xgb = GridSearchCV(XGBClassifier(), param_grid, refit = True, verbose = 3,
    ↪cv = 10, n_jobs = -1)

# fitting the model for grid search
grid_xgb.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_xgb.best_params_)

# print how our model looks after hyper-parameter tuning
```

```
print(grid_xgb.best_estimator_)
print(grid_xgb.best_score_)
```

Fitting 10 folds for each of 15 candidates, totalling 150 fits

```
{'gamma': 0.01, 'n_estimators': 150}
```

```
XGBClassifier(base_score=0.5, booster='gbtree', callbacks=None,
              colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,
              early_stopping_rounds=None, enable_categorical=False,
              eval_metric=None, gamma=0.01, gpu_id=-1, grow_policy='depthwise',
              importance_type=None, interaction_constraints='',
              learning_rate=0.300000012, max_bin=256, max_cat_to_onehot=4,
              max_delta_step=0, max_depth=6, max_leaves=0, min_child_weight=1,
              missing=nan, monotone_constraints='()', n_estimators=150,
              n_jobs=0, num_parallel_tree=1, predictor='auto', random_state=0,
              reg_alpha=0, reg_lambda=1, ...)
```

0.9642167578334326

```
[ ]: xgb_model = grid_xgb.best_estimator_
      #xgb_model = xgb.fit(X_train,y_train)
```

```
[ ]: xgb_predict=xgb_model.predict(X_test)
```


```
[ ]: print('The accuracy of XGBoost Classifier is: ' , 100.0 *
      ↪accuracy_score(xgb_predict,y_test))
```

The accuracy of XGBoost Classifier is: 96.65461121157324

```
[ ]: print(classification_report(y_test, xgb_predict))
```

	precision	recall	f1-score	support
0	0.97	0.95	0.96	472
1	0.97	0.98	0.97	634
accuracy			0.97	1106
macro avg	0.97	0.96	0.97	1106
weighted avg	0.97	0.97	0.97	1106

```
[ ]: sns.heatmap(confusion_matrix(y_test, xgb_predict), annot=True, fmt='g',
      ↪cmap='Blues')
plt.title("XGBClassifier")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
```



```
[ ]: from sklearn.ensemble import GradientBoostingClassifier

# defining parameter range
param_grid = {
    "learning_rate": [.1,.5,1],
    "n_estimators": [50,100,150,200,250]
}

grid_gbc = GridSearchCV(GradientBoostingClassifier(), param_grid, refit = True,
    verbose = 3, cv = 10, n_jobs = -1)

# fitting the model for grid search
grid_gbc.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_gbc.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_gbc.best_estimator_)
```

```
print(grid_gbc.best_score_)
```

Fitting 10 folds for each of 15 candidates, totalling 150 fits  
{'learning\_rate': 1, 'n\_estimators': 250}  
GradientBoostingClassifier(learning\_rate=1, n\_estimators=250)  
0.9622069098005117

```
[ ]: gbc_model = grid_gbc.best_estimator_  
      #gbc_model = gbc.fit(X_train,y_train.values.ravel())  
  
      #clf = GradientBoostingClassifier(n_estimators=100, learning_rate=1.0,  
      #    max_depth=1, random_state=0).fit(X_train, y_train)  
      #clf.score(X_test, y_test)
```

```
[ ]: gbc_predict = gbc_model.predict(X_test)
```


```
[ ]: print('The accuracy of GradientBoost Classifier is: ', 100.0 *  
      ↪accuracy_score(gbc_predict,y_test))
```

The accuracy of GradientBoost Classifier is: 96.74502712477397

```
[ ]: print(classification_report(y_test, gbc_predict))
```

	precision	recall	f1-score	support
0	0.97	0.96	0.96	472
1	0.97	0.97	0.97	634
accuracy			0.97	1106
macro avg	0.97	0.97	0.97	1106
weighted avg	0.97	0.97	0.97	1106

```
[ ]: sns.heatmap(confusion_matrix(y_test, gbc_predict), annot=True, fmt='g',  
      ↪cmap='Blues')  
plt.title("GradientBoostingClassifier")  
plt.xlabel('Predicted Class')  
plt.ylabel('Original Class')  
plt.show()
```



```
[ ]: # gbc_model.get_params().keys()
```

```
[ ]: # import inspect
# import sklearn
# import xgboost

# models = [xgboost.XGBClassifier]
# for m in models:
#     hyperparams = inspect.signature(m.__init__)
#     print(hyperparams)
# #or
# xgb_model.get_params().keys()
```

```
[ ]: from sklearn.ensemble import BaggingClassifier
from sklearn.tree import DecisionTreeClassifier

# defining parameter range
param_grid = {
    "base_estimator": [DecisionTreeClassifier()],
    "n_estimators": [50,100,150,200,250]
```

```

}

grid_bag = GridSearchCV(BaggingClassifier(), param_grid, refit = True, verbose_
    ↪= 3, cv = 10, n_jobs = -1)

# fitting the model for grid search
grid_bag.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_bag.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_bag.best_estimator_)
print(grid_bag.best_score_)

```

```

Fitting 10 folds for each of 5 candidates, totalling 50 fits
{'base_estimator': DecisionTreeClassifier(), 'n_estimators': 200}
BaggingClassifier(base_estimator=DecisionTreeClassifier(), n_estimators=200)
0.9625088217748703

```

```

[ ]: bag_model = grid_bag.best_estimator_
    ↪bag_model = bag.fit(X_train, y_train.values.ravel())

```

```

[ ]: bag_predict = bag_model.predict(X_test)

```

```

[ ]: print('The accuracy of Bagging Classifier is: ', 100.0 *_
    ↪accuracy_score(y_test, bag_predict))

```

```

The accuracy of Bagging Classifier is: 96.74502712477397

```

```

[ ]: print(classification_report(y_test, bag_predict))

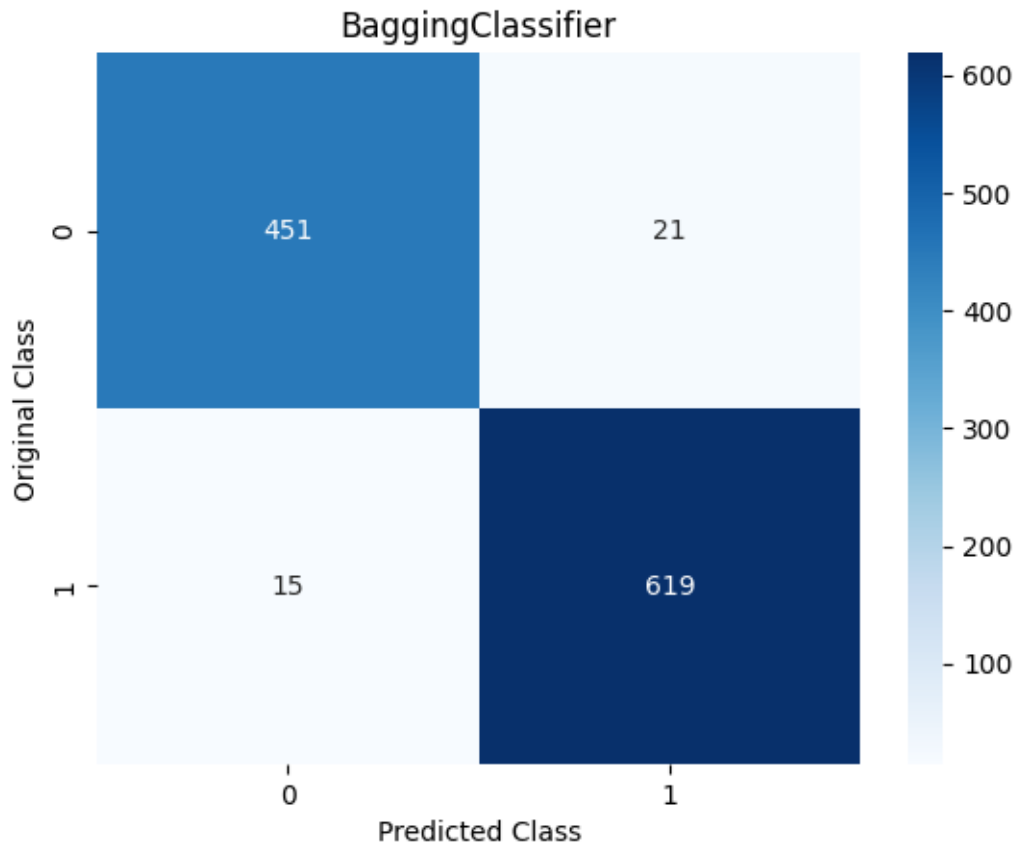
```

	precision	recall	f1-score	support
0	0.97	0.96	0.96	472
1	0.97	0.98	0.97	634
accuracy			0.97	1106
macro avg	0.97	0.97	0.97	1106
weighted avg	0.97	0.97	0.97	1106

```

[ ]: sns.heatmap(confusion_matrix(y_test, bag_predict), annot=True, fmt='g',_
    ↪cmap='Blues')
plt.title("BaggingClassifier")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()

```



```
[ ]: from sklearn.ensemble import RandomForestClassifier

# defining parameter range
param_grid = {
    "n_estimators": [50,100,150,200,250]
}

grid_rfc = GridSearchCV(RandomForestClassifier(), param_grid, refit = True,
    verbose = 3, cv = 10, n_jobs = -1)

# fitting the model for grid search
grid_rfc.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_rfc.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_rfc.best_estimator_)
print(grid_rfc.best_score_)
```

```
Fitting 10 folds for each of 5 candidates, totalling 50 fits
{'n_estimators': 50}
RandomForestClassifier(n_estimators=50)
0.9643175636734982
```

```
[ ]: rfc_model = grid_rfc.best_estimator_
      #rfc_model = rfc.fit(X_train,y_train.values.ravel())
```

```
[ ]: rfc_predict = rfc_model.predict(X_test)
```

```
[ ]: print('The accuracy of RandomForest Classifier is: ' , 100.0 *
      ↪accuracy_score(rfc_predict,y_test))
```


The accuracy of RandomForest Classifier is: 96.83544303797468

```
[ ]: print(classification_report(y_test, rfc_predict))
```


	precision	recall	f1-score	support
0	0.97	0.96	0.96	472
1	0.97	0.98	0.97	634
accuracy			0.97	1106
macro avg	0.97	0.97	0.97	1106
weighted avg	0.97	0.97	0.97	1106

```
[ ]: sns.heatmap(confusion_matrix(y_test, rfc_predict), annot=True, fmt='g',
      ↪cmap='Blues')
plt.title("RandomForestClassifier")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
```





```
[ ]: estimators =  
    ↳ [logr_model,knn_model,svc_model,nusvc_model,lsvc_model,xgb_model,ada_model,gbc_model,bag_mo  
  
for estimator in estimators:  
    RocCurveDisplay.from_estimator(estimator,X_test,y_test,ax=plt.gca())
```



```
[ ]: import tensorflow as tf
#from tensorflow.keras.datasets import imdb
from keras.layers import Embedding, Dense, LSTM, BatchNormalization
from keras.losses import BinaryCrossentropy
from keras.models import Sequential
from keras.optimizers import Adam
#from tensorflow.keras.preprocessing.sequence import pad_sequences

# Model configuration
additional_metrics = ['accuracy']
batch_size = 32
#embedding_output_dims = (X_train.shape[1])
loss_function = BinaryCrossentropy()
#max_sequence_length = (X_train.shape[1])
#num_distinct_words = (X_train.shape[1])
number_of_epochs = 100
optimizer = Adam()
validation_split = 0.20
verbosity_mode = 1

# reshape from [samples, features] into [samples, timesteps, features]
```

```

timesteps = 1
X_train_reshape = X_train.values.ravel().reshape(X_train.shape[0],timesteps,
↳X_train.shape[1])
X_test_reshape = X_test.values.ravel().reshape(X_test.shape[0],timesteps,
↳X_test.shape[1])

# Disable eager execution
#tf.compat.v1.disable_eager_execution()

# Load dataset
# (x_train, y_train), (x_test, y_test) = imdb.
↳load_data(num_words=num_distinct_words)
# print(x_train.shape)
# print(x_test.shape)

# Pad all sequences
# padded_inputs = pad_sequences(X_train, maxlen=max_sequence_length, value = 0.
↳0) # 0.0 because it corresponds with <PAD>
# padded_inputs_test = pad_sequences(X_test, maxlen=max_sequence_length, value
↳= 0.0) # 0.0 because it corresponds with <PAD>

# Define the Keras model
def build_model_lstm():
    model = Sequential()
    #model.add(Embedding(num_distinct_words, embedding_output_dims,
↳input_length=max_sequence_length))
    model.add(LSTM(100, input_shape = (timesteps,X_train_reshape.shape[2])))
    model.add(BatchNormalization())
    model.add(Dense(50, activation='relu'))
    model.add(Dense(25, activation='relu'))
    model.add(Dense(10, activation='relu'))
    model.add(Dense(1, activation='sigmoid'))

    # Compile the model
    model.compile(optimizer=optimizer, loss=loss_function,
↳metrics=additional_metrics)
    return model

#from keras.wrappers.scikit_learn import KerasClassifier
lstm_model = build_model_lstm()
# Give a summary
lstm_model.summary()

# Train the model

```

```

history = lstm_model.fit(X_train_reshape, y_train.values.ravel(),
    ↪batch_size=batch_size, epochs=number_of_epochs, verbose=verbosity_mode,
    ↪validation_split=validation_split)

# Test the model after training
#lstm_predict = lstm_model.predict(X_test_reshape)
test_results = lstm_model.evaluate(X_test_reshape, y_test.values.ravel(),
    ↪verbose=False)
print(f'Test results - Loss: {test_results[0]} - Accuracy:
    ↪{100*test_results[1]}%')

```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
lstm_2 (LSTM)	(None, 100)	48000
batch_normalization_2 (Batch Normalization)	(None, 100)	400
dense_8 (Dense)	(None, 50)	5050
dense_9 (Dense)	(None, 25)	1275
Layer (type)	Output Shape	Param #
lstm_2 (LSTM)	(None, 100)	48000
batch_normalization_2 (Batch Normalization)	(None, 100)	400
dense_8 (Dense)	(None, 50)	5050
dense_9 (Dense)	(None, 25)	1275
dense_10 (Dense)	(None, 10)	260
dense_11 (Dense)	(None, 1)	11

Total params: 54,996  
 Trainable params: 54,796  
 Non-trainable params: 200

Epoch 1/100  
 249/249 [=====] - 4s 6ms/step - loss: 0.2307 -  
 accuracy: 0.9118 - val\_loss: 0.3389 - val\_accuracy: 0.9387

Epoch 2/100  
249/249 [=====] - 1s 5ms/step - loss: 0.1739 -  
accuracy: 0.9273 - val\_loss: 0.1605 - val\_accuracy: 0.9447  
Epoch 3/100  
249/249 [=====] - 1s 5ms/step - loss: 0.1549 -  
accuracy: 0.9358 - val\_loss: 0.1293 - val\_accuracy: 0.9437  
Epoch 4/100  
249/249 [=====] - 1s 4ms/step - loss: 0.1445 -  
accuracy: 0.9382 - val\_loss: 0.1344 - val\_accuracy: 0.9447  
Epoch 5/100  
249/249 [=====] - 1s 4ms/step - loss: 0.1371 -  
accuracy: 0.9422 - val\_loss: 0.1220 - val\_accuracy: 0.9482  
Epoch 6/100  
249/249 [=====] - 1s 4ms/step - loss: 0.1275 -  
accuracy: 0.9460 - val\_loss: 0.1169 - val\_accuracy: 0.9503  
Epoch 7/100  
249/249 [=====] - 1s 4ms/step - loss: 0.1224 -  
accuracy: 0.9489 - val\_loss: 0.1202 - val\_accuracy: 0.9503  
Epoch 8/100  
249/249 [=====] - 1s 4ms/step - loss: 0.1182 -  
accuracy: 0.9506 - val\_loss: 0.1154 - val\_accuracy: 0.9543  
Epoch 9/100  
249/249 [=====] - 1s 4ms/step - loss: 0.1150 -  
accuracy: 0.9506 - val\_loss: 0.1184 - val\_accuracy: 0.9523  
Epoch 10/100  
249/249 [=====] - 1s 4ms/step - loss: 0.1095 -  
accuracy: 0.9556 - val\_loss: 0.1120 - val\_accuracy: 0.9528  
Epoch 11/100  
249/249 [=====] - 1s 4ms/step - loss: 0.1041 -  
accuracy: 0.9554 - val\_loss: 0.1174 - val\_accuracy: 0.9533  
Epoch 12/100  
249/249 [=====] - 1s 4ms/step - loss: 0.1027 -  
accuracy: 0.9570 - val\_loss: 0.1086 - val\_accuracy: 0.9523  
Epoch 13/100  
249/249 [=====] - 1s 5ms/step - loss: 0.1041 -  
accuracy: 0.9567 - val\_loss: 0.1154 - val\_accuracy: 0.9457  
Epoch 14/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0986 -  
accuracy: 0.9588 - val\_loss: 0.1106 - val\_accuracy: 0.9553  
Epoch 15/100  
249/249 [=====] - 1s 5ms/step - loss: 0.0958 -  
accuracy: 0.9593 - val\_loss: 0.1135 - val\_accuracy: 0.9533  
Epoch 16/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0941 -  
accuracy: 0.9603 - val\_loss: 0.1097 - val\_accuracy: 0.9563  
Epoch 17/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0936 -  
accuracy: 0.9587 - val\_loss: 0.1126 - val\_accuracy: 0.9563

Epoch 18/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0885 -  
accuracy: 0.9619 - val\_loss: 0.1025 - val\_accuracy: 0.9583  
Epoch 19/100  
249/249 [=====] - 1s 5ms/step - loss: 0.0875 -  
accuracy: 0.9621 - val\_loss: 0.1054 - val\_accuracy: 0.9568  
Epoch 20/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0850 -  
accuracy: 0.9619 - val\_loss: 0.0981 - val\_accuracy: 0.9608  
Epoch 21/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0841 -  
accuracy: 0.9642 - val\_loss: 0.1072 - val\_accuracy: 0.9568  
Epoch 22/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0812 -  
accuracy: 0.9644 - val\_loss: 0.1203 - val\_accuracy: 0.9573  
Epoch 23/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0810 -  
accuracy: 0.9663 - val\_loss: 0.1092 - val\_accuracy: 0.9583  
Epoch 24/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0808 -  
accuracy: 0.9646 - val\_loss: 0.1190 - val\_accuracy: 0.9563  
Epoch 25/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0783 -  
accuracy: 0.9656 - val\_loss: 0.1120 - val\_accuracy: 0.9608  
Epoch 26/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0750 -  
accuracy: 0.9667 - val\_loss: 0.1068 - val\_accuracy: 0.9613  
Epoch 27/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0730 -  
accuracy: 0.9673 - val\_loss: 0.1080 - val\_accuracy: 0.9578  
Epoch 28/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0724 -  
accuracy: 0.9681 - val\_loss: 0.1273 - val\_accuracy: 0.9523  
Epoch 29/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0702 -  
accuracy: 0.9683 - val\_loss: 0.1055 - val\_accuracy: 0.9583  
Epoch 30/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0699 -  
accuracy: 0.9685 - val\_loss: 0.1160 - val\_accuracy: 0.9593  
Epoch 31/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0693 -  
accuracy: 0.9690 - val\_loss: 0.1142 - val\_accuracy: 0.9578  
Epoch 32/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0695 -  
accuracy: 0.9706 - val\_loss: 0.1146 - val\_accuracy: 0.9548  
Epoch 33/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0695 -  
accuracy: 0.9696 - val\_loss: 0.1130 - val\_accuracy: 0.9618

Epoch 34/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0672 -  
accuracy: 0.9703 - val\_loss: 0.1137 - val\_accuracy: 0.9608  
Epoch 35/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0658 -  
accuracy: 0.9707 - val\_loss: 0.1183 - val\_accuracy: 0.9563  
Epoch 36/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0683 -  
accuracy: 0.9686 - val\_loss: 0.1067 - val\_accuracy: 0.9613  
Epoch 37/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0658 -  
accuracy: 0.9714 - val\_loss: 0.1128 - val\_accuracy: 0.9583  
Epoch 38/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0682 -  
accuracy: 0.9702 - val\_loss: 0.1305 - val\_accuracy: 0.9538  
Epoch 39/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0692 -  
accuracy: 0.9697 - val\_loss: 0.1096 - val\_accuracy: 0.9603  
Epoch 40/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0624 -  
accuracy: 0.9712 - val\_loss: 0.1094 - val\_accuracy: 0.9553  
Epoch 41/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0659 -  
accuracy: 0.9714 - val\_loss: 0.1091 - val\_accuracy: 0.9603  
Epoch 42/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0606 -  
accuracy: 0.9726 - val\_loss: 0.1157 - val\_accuracy: 0.9608  
Epoch 43/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0617 -  
accuracy: 0.9736 - val\_loss: 0.1210 - val\_accuracy: 0.9588  
Epoch 44/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0595 -  
accuracy: 0.9732 - val\_loss: 0.1351 - val\_accuracy: 0.9578  
Epoch 45/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0626 -  
accuracy: 0.9725 - val\_loss: 0.1234 - val\_accuracy: 0.9608  
Epoch 46/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0630 -  
accuracy: 0.9720 - val\_loss: 0.1094 - val\_accuracy: 0.9608  
Epoch 47/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0595 -  
accuracy: 0.9744 - val\_loss: 0.1094 - val\_accuracy: 0.9578  
Epoch 48/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0582 -  
accuracy: 0.9735 - val\_loss: 0.1154 - val\_accuracy: 0.9568  
Epoch 49/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0598 -  
accuracy: 0.9732 - val\_loss: 0.1136 - val\_accuracy: 0.9623

Epoch 50/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0588 -  
accuracy: 0.9732 - val\_loss: 0.1121 - val\_accuracy: 0.9648  
Epoch 51/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0599 -  
accuracy: 0.9734 - val\_loss: 0.1124 - val\_accuracy: 0.9618  
Epoch 52/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0561 -  
accuracy: 0.9735 - val\_loss: 0.1165 - val\_accuracy: 0.9608  
Epoch 53/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0610 -  
accuracy: 0.9742 - val\_loss: 0.1159 - val\_accuracy: 0.9613  
Epoch 54/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0556 -  
accuracy: 0.9761 - val\_loss: 0.1159 - val\_accuracy: 0.9593  
Epoch 55/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0578 -  
accuracy: 0.9744 - val\_loss: 0.1069 - val\_accuracy: 0.9613  
Epoch 56/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0547 -  
accuracy: 0.9745 - val\_loss: 0.1092 - val\_accuracy: 0.9598  
Epoch 57/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0536 -  
accuracy: 0.9758 - val\_loss: 0.1223 - val\_accuracy: 0.9608  
Epoch 58/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0506 -  
accuracy: 0.9771 - val\_loss: 0.1267 - val\_accuracy: 0.9628  
Epoch 59/100  
249/249 [=====] - 1s 5ms/step - loss: 0.0522 -  
accuracy: 0.9764 - val\_loss: 0.1220 - val\_accuracy: 0.9663  
Epoch 60/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0519 -  
accuracy: 0.9774 - val\_loss: 0.1112 - val\_accuracy: 0.9623  
Epoch 61/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0519 -  
accuracy: 0.9769 - val\_loss: 0.1375 - val\_accuracy: 0.9623  
Epoch 62/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0581 -  
accuracy: 0.9742 - val\_loss: 0.1410 - val\_accuracy: 0.9578  
Epoch 63/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0561 -  
accuracy: 0.9735 - val\_loss: 0.1165 - val\_accuracy: 0.9613  
Epoch 64/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0543 -  
accuracy: 0.9763 - val\_loss: 0.1278 - val\_accuracy: 0.9608  
Epoch 65/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0523 -  
accuracy: 0.9744 - val\_loss: 0.1167 - val\_accuracy: 0.9578



Epoch 66/100  
249/249 [=====] - 1s 5ms/step - loss: 0.0516 -  
accuracy: 0.9783 - val\_loss: 0.1178 - val\_accuracy: 0.9608  
Epoch 67/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0492 -  
accuracy: 0.9776 - val\_loss: 0.1278 - val\_accuracy: 0.9618  
Epoch 68/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0496 -  
accuracy: 0.9770 - val\_loss: 0.1227 - val\_accuracy: 0.9633  
Epoch 69/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0554 -  
accuracy: 0.9746 - val\_loss: 0.1302 - val\_accuracy: 0.9648  
Epoch 70/100  
249/249 [=====] - 1s 5ms/step - loss: 0.0534 -  
accuracy: 0.9759 - val\_loss: 0.1109 - val\_accuracy: 0.9653  
Epoch 71/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0505 -  
accuracy: 0.9764 - val\_loss: 0.1197 - val\_accuracy: 0.9623  
Epoch 72/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0497 -  
accuracy: 0.9760 - val\_loss: 0.1224 - val\_accuracy: 0.9628  
Epoch 73/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0518 -  
accuracy: 0.9773 - val\_loss: 0.1334 - val\_accuracy: 0.9603  
Epoch 74/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0507 -  
accuracy: 0.9766 - val\_loss: 0.1117 - val\_accuracy: 0.9638  
Epoch 75/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0475 -  
accuracy: 0.9776 - val\_loss: 0.1209 - val\_accuracy: 0.9643  
Epoch 76/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0493 -  
accuracy: 0.9785 - val\_loss: 0.1409 - val\_accuracy: 0.9618  
Epoch 77/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0486 -  
accuracy: 0.9770 - val\_loss: 0.1264 - val\_accuracy: 0.9623  
Epoch 78/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0502 -  
accuracy: 0.9763 - val\_loss: 0.1176 - val\_accuracy: 0.9608  
Epoch 79/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0466 -  
accuracy: 0.9773 - val\_loss: 0.1271 - val\_accuracy: 0.9593  
Epoch 80/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0486 -  
accuracy: 0.9765 - val\_loss: 0.1248 - val\_accuracy: 0.9638  
Epoch 81/100  
249/249 [=====] - 1s 5ms/step - loss: 0.0508 -  
accuracy: 0.9773 - val\_loss: 0.1257 - val\_accuracy: 0.9633

Epoch 82/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0501 -  
accuracy: 0.9763 - val\_loss: 0.1236 - val\_accuracy: 0.9613

Epoch 83/100  
249/249 [=====] - 1s 5ms/step - loss: 0.0482 -  
accuracy: 0.9773 - val\_loss: 0.1340 - val\_accuracy: 0.9613

Epoch 84/100  
249/249 [=====] - 1s 5ms/step - loss: 0.0468 -  
accuracy: 0.9783 - val\_loss: 0.1389 - val\_accuracy: 0.9593

Epoch 85/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0501 -  
accuracy: 0.9768 - val\_loss: 0.1433 - val\_accuracy: 0.9583

Epoch 86/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0463 -  
accuracy: 0.9776 - val\_loss: 0.1319 - val\_accuracy: 0.9628

Epoch 87/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0435 -  
accuracy: 0.9799 - val\_loss: 0.1170 - val\_accuracy: 0.9593

Epoch 88/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0444 -  
accuracy: 0.9799 - val\_loss: 0.1281 - val\_accuracy: 0.9598

Epoch 89/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0443 -  
accuracy: 0.9785 - val\_loss: 0.1297 - val\_accuracy: 0.9623

Epoch 90/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0441 -  
accuracy: 0.9781 - val\_loss: 0.1321 - val\_accuracy: 0.9618

Epoch 91/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0520 -  
accuracy: 0.9755 - val\_loss: 0.1391 - val\_accuracy: 0.9633

Epoch 92/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0413 -  
accuracy: 0.9803 - val\_loss: 0.1518 - val\_accuracy: 0.9608

Epoch 93/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0440 -  
accuracy: 0.9778 - val\_loss: 0.1303 - val\_accuracy: 0.9638

Epoch 94/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0477 -  
accuracy: 0.9768 - val\_loss: 0.1466 - val\_accuracy: 0.9658

Epoch 95/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0438 -  
accuracy: 0.9786 - val\_loss: 0.1415 - val\_accuracy: 0.9628

Epoch 96/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0535 -  
accuracy: 0.9768 - val\_loss: 0.1520 - val\_accuracy: 0.9588

Epoch 97/100  
249/249 [=====] - 1s 4ms/step - loss: 0.0482 -  
accuracy: 0.9774 - val\_loss: 0.1477 - val\_accuracy: 0.9608


```
Epoch 98/100
249/249 [=====] - 1s 4ms/step - loss: 0.0482 -
accuracy: 0.9788 - val_loss: 0.1267 - val_accuracy: 0.9628
Epoch 99/100
249/249 [=====] - 1s 4ms/step - loss: 0.0471 -
accuracy: 0.9774 - val_loss: 0.1453 - val_accuracy: 0.9608
Epoch 100/100
249/249 [=====] - 1s 4ms/step - loss: 0.0425 -
accuracy: 0.9798 - val_loss: 0.1459 - val_accuracy: 0.9643
Test results - Loss: 0.1807340830564499 - Accuracy: 96.02169990539551%
```

```
[ ]: lstm_predict_proba = lstm_model.predict(X_test_reshape, batch_size=32)
lstm_predict_class = (lstm_predict_proba > 0.5).astype("int32")
print(classification_report(y_test, lstm_predict_class))
```

```
35/35 [=====] - 1s 2ms/step
```

	precision	recall	f1-score	support
0	0.96	0.94	0.95	472
1	0.96	0.97	0.97	634
accuracy			0.96	1106
macro avg	0.96	0.96	0.96	1106
weighted avg	0.96	0.96	0.96	1106

```
[ ]: sns.heatmap(confusion_matrix(y_test, lstm_predict_class), annot=True, fmt='g',
cmap='Blues')
plt.title("LSTM")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
```



```
[ ]: RocCurveDisplay.from_predictions(y_test,lstm_predict_class)
plt.show()
```

