

## chi\_sq\_87 8020 split .05 threshold

January 3, 2023

```
[ ]: # Importing the packages
import sys
import numpy as np
np.set_printoptions(threshold=sys.maxsize)
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
import sklearn
import random
from sklearn.metrics import ↵
    ↵confusion_matrix, accuracy_score, classification_report, RocCurveDisplay, ConfusionMatrixDisplay

[ ]: pd.set_option('display.max_rows', None)
pd.set_option('display.max_columns', None)
pd.set_option('display.width', None)
pd.set_option('display.max_colwidth', None)

[ ]: # Importing the dataset
df = pd.read_csv('dataset_phishing.csv')
df.drop(['url'], axis=1, inplace=True)
#df.head(50)

[ ]: # if your dataset contains missing value, check which column has missing values
#df.isnull().sum()

[ ]: #df.dropna(inplace=True)

[ ]: from sklearn import preprocessing

col = [df.columns[-1]]

lab_en= preprocessing.LabelEncoder()

for c in col:
    df[c]= lab_en.fit_transform(df[c])

#df.head(50)
```

```
[ ]: a=len(df[df.status==0])
      b=len(df[df.status==1])
```

```
[ ]: print("Count of Legitimate Websites = ", a)
      print("Count of Phishy Websites = ", b)
```

Count of Legitimate Websites = 5715  
Count of Phishy Websites = 5715

```
[ ]: X = df.drop(['status'], axis=1, inplace=False)
      #X.head()
      #same work
      ##inplace true modifies the og data & does not return anything
      ##inplace false does not modify og data but returns something which we store in
      ↪ a var
      # X= df.drop(columns='Result')
      # X.head()
```

```
[ ]: #df.head()
```

```
[ ]: y = df['status']
      y = pd.DataFrame(y)
      y.head()
```

```
[ ]:      status
      0      0
      1      1
      2      1
      3      0
      4      0
```

```
[ ]: # separate dataset into train and test
      from cProfile import label
      from sklearn.model_selection import train_test_split
      X_train, X_test, y_train, y_test = train_test_split(
          X,
          y,
          test_size=0.2,
          random_state=10)

      X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
[ ]: ((9144, 87), (2286, 87), (9144, 1), (2286, 1))
```

```
[ ]: #X_test.head()
```

```
[ ]: from sklearn.preprocessing import MinMaxScaler
```

```

scaler= MinMaxScaler()

col_X_train = [X_train.columns[:]]

for c in col_X_train:
    X_train[c]= scaler.fit_transform(X_train[c])

#X_train.head(5)

```

```

[ ]: col_X_test = [X_test.columns[:]]

for c in col_X_test:
    X_test[c]= scaler.transform(X_test[c])

#X_test.head(5)

```

```

[ ]: #perform chi square test
from sklearn.feature_selection import chi2
f_p_values = chi2(X_train,y_train)

```

```

[ ]: f_p_values

```

```

[ ]: (array([2.23370641e+01, 1.69566642e+01, 7.92634978e+02, 2.14222888e+01,
            8.35897474e+00, 5.30407662e+01, 2.53295162e+02, 6.10286729e+01,
            nan, 9.21347603e+01, 2.48852581e+00, 5.32938701e+00,
            1.36127244e+00, 2.76337058e+01, 6.97249509e+00, 3.02335502e+01,
            3.37730377e-02, 2.83381683e+01, 3.15422397e+00, 1.94038283e-02,
            5.01524054e+02, 4.11373688e+01, 4.68200813e+01, 2.09854284e+01,
            4.59133775e+01, 2.65045377e+02, 1.94693671e+02, 2.98821218e+00,
            1.97645536e+00, 5.07490559e+01, 3.59432236e+02, 1.31736523e+02,
            1.89723928e+01, 3.31523214e+02, 1.71903269e+00, 8.25729068e+01,
            7.75004553e-06, 8.27679909e-01, 2.98821218e+01, 1.99134491e+01,
            9.70793366e-02, 9.44986388e-01, 4.73779494e+01, 4.67727391e+00,
            2.67052777e+01, 6.69644804e+00, 4.14067508e+01, 8.24823685e+00,
            1.81658385e+01, 1.75476572e+01, 2.20716199e+02, 7.15055802e+01,
            4.28310413e+01, 3.52188528e+01, 1.13779295e+02, 1.73543170e+02,
            7.52046451e+01, 1.25671887e+02, 2.92566949e+01, nan,
            4.45199645e+00, nan, 4.22631863e+01, nan,
            3.50467989e+00, 2.62448221e+00, 1.10884273e+02, 1.01180269e+02,
            nan, 1.71196163e+02, 1.05123590e+02, nan,
            1.34913309e+00, 3.73506517e+01, 1.17153886e+02, 1.15083440e-01,
            7.30367272e-02, 3.23965734e+02, 2.40907086e+02, 1.55570635e+02,
            3.98630823e+01, 1.01301038e+01, 1.86177846e+02, 1.26203754e+01,
            1.30175372e+02, 2.24891968e+03, 4.76130808e+02]),
      array([2.28748708e-006, 3.82428387e-005, 2.15456012e-174, 3.68462862e-006,
            3.83787046e-003, 3.26697163e-013, 4.96687374e-057, 5.62495784e-015,
            nan, 8.09705637e-022, 1.14679097e-001, 2.09687588e-002,

```

```

2.43317079e-001, 1.46601796e-007, 8.27720043e-003, 3.83026070e-008,
8.54190424e-001, 1.01867512e-007, 7.57306913e-002, 8.89214967e-001,
4.42963423e-111, 1.41897299e-010, 7.78119832e-012, 4.62789878e-006,
1.23599050e-011, 1.36390512e-059, 3.00516293e-044, 8.38727207e-002,
1.59764326e-001, 1.04963046e-012, 3.74294379e-080, 1.70855173e-030,
1.32623548e-005, 4.47745841e-074, 1.89817606e-001, 1.01843501e-019,
9.97778780e-001, 3.62944266e-001, 4.59126916e-008, 8.10283030e-006,
7.55363102e-001, 3.30998761e-001, 5.85374855e-012, 3.05642145e-002,
2.36973375e-007, 9.66051752e-003, 1.23629710e-010, 4.07916034e-003,
2.02478406e-005, 2.80195804e-005, 6.31194922e-050, 2.76475119e-017,
5.96776479e-011, 2.94657202e-009, 1.45650442e-026, 1.24551325e-039,
4.24367530e-018, 3.62767027e-029, 6.33968046e-008, nan,
3.48604695e-002, nan, 7.97802714e-011, nan,
6.11956709e-002, 1.05226826e-001, 6.27276739e-026, 8.39817380e-024,
nan, 4.05431738e-039, 1.14750659e-024, nan,
2.45429728e-001, 9.86878980e-010, 2.65636383e-027, 7.34429313e-001,
7.86965560e-001, 1.98196025e-072, 2.49414539e-054, 1.05073706e-035,
2.72403538e-010, 1.45867197e-003, 2.17104281e-042, 3.81564706e-004,
3.75126966e-030, 0.00000000e+000, 1.48483498e-105]))

```

```

[ ]: #The less the p_values the more important that feature is
p_values = pd.Series(f_p_values[1])
p_values.index = X_train.columns
p_values

```

```

[ ]: length_url          2.287487e-06
length_hostname        3.824284e-05
ip                    2.154560e-174
nb_dots                3.684629e-06
nb_hyphens             3.837870e-03
nb_at                  3.266972e-13
nb_qm                  4.966874e-57
nb_and                 5.624958e-15
nb_or                  NaN
nb_eq                  8.097056e-22
nb_underscore          1.146791e-01
nb_tilde               2.096876e-02
nb_percent             2.433171e-01
nb_slash               1.466018e-07
nb_star                8.277200e-03
nb_colon               3.830261e-08
nb_comma               8.541904e-01
nb_semicolumn          1.018675e-07
nb_dollar              7.573069e-02
nb_space               8.892150e-01
nb_www                 4.429634e-111
nb_com                 1.418973e-10

```

nb_dslash	7.781198e-12
http_in_path	4.627899e-06
https_token	1.235990e-11
ratio_digits_url	1.363905e-59
ratio_digits_host	3.005163e-44
punycode	8.387272e-02
port	1.597643e-01
tld_in_path	1.049630e-12
tld_in_subdomain	3.742944e-80
abnormal_subdomain	1.708552e-30
nb_subdomains	1.326235e-05
prefix_suffix	4.477458e-74
random_domain	1.898176e-01
shortening_service	1.018435e-19
path_extension	9.977788e-01
nb_redirection	3.629443e-01
nb_external_redirection	4.591269e-08
length_words_raw	8.102830e-06
char_repeat	7.553631e-01
shortest_words_raw	3.309988e-01
shortest_word_host	5.853749e-12
shortest_word_path	3.056421e-02
longest_words_raw	2.369734e-07
longest_word_host	9.660518e-03
longest_word_path	1.236297e-10
avg_words_raw	4.079160e-03
avg_word_host	2.024784e-05
avg_word_path	2.801958e-05
phish_hints	6.311949e-50
domain_in_brand	2.764751e-17
brand_in_subdomain	5.967765e-11
brand_in_path	2.946572e-09
suspicious_tld	1.456504e-26
statistical_report	1.245513e-39
nb_hyperlinks	4.243675e-18
ratio_intHyperlinks	3.627670e-29
ratio_extHyperlinks	6.339680e-08
ratio_nullHyperlinks	NaN
nb_extCSS	3.486047e-02
ratio_intRedirection	NaN
ratio_extRedirection	7.978027e-11
ratio_intErrors	NaN
ratio_extErrors	6.119567e-02
login_form	1.052268e-01
external_favicon	6.272767e-26
links_in_tags	8.398174e-24
submit_email	NaN

ratio_intMedia	4.054317e-39
ratio_extMedia	1.147507e-24
sfh	NaN
iframe	2.454297e-01
popup_window	9.868790e-10
safe_anchor	2.656364e-27
onmouseover	7.344293e-01
right_clic	7.869656e-01
empty_title	1.981960e-72
domain_in_title	2.494145e-54
domain_with_copyright	1.050737e-35
whois_registered_domain	2.724035e-10
domain_registration_length	1.458672e-03
domain_age	2.171043e-42
web_traffic	3.815647e-04
dns_record	3.751270e-30
google_index	0.000000e+00
page_rank	1.484835e-105
dtype:	float64

```
[ ]: #sort p_values to check which feature has the lowest values
p_values = p_values.sort_values(ascending = False)
p_values
```

[ ]: path_extension	9.977788e-01
nb_space	8.892150e-01
nb_comma	8.541904e-01
right_clic	7.869656e-01
char_repeat	7.553631e-01
onmouseover	7.344293e-01
nb_redirection	3.629443e-01
shortest_words_raw	3.309988e-01
iframe	2.454297e-01
nb_percent	2.433171e-01
random_domain	1.898176e-01
port	1.597643e-01
nb_underscore	1.146791e-01
login_form	1.052268e-01
punycode	8.387272e-02
nb_dollar	7.573069e-02
ratio_extErrors	6.119567e-02
nb_extCSS	3.486047e-02
shortest_word_path	3.056421e-02
nb_tilde	2.096876e-02
longest_word_host	9.660518e-03
nb_star	8.277200e-03
avg_words_raw	4.079160e-03

nb_hyphens	3.837870e-03
domain_registration_length	1.458672e-03
web_traffic	3.815647e-04
length_hostname	3.824284e-05
avg_word_path	2.801958e-05
avg_word_host	2.024784e-05
nb_subdomains	1.326235e-05
length_words_raw	8.102830e-06
http_in_path	4.627899e-06
nb_dots	3.684629e-06
length_url	2.287487e-06
longest_words_raw	2.369734e-07
nb_slash	1.466018e-07
nb_semicolumn	1.018675e-07
ratio_extHyperlinks	6.339680e-08
nb_external_redirection	4.591269e-08
nb_colon	3.830261e-08
brand_in_path	2.946572e-09
popup_window	9.868790e-10
whois_registered_domain	2.724035e-10
nb_com	1.418973e-10
longest_word_path	1.236297e-10
ratio_extRedirection	7.978027e-11
brand_in_subdomain	5.967765e-11
https_token	1.235990e-11
nb_dslash	7.781198e-12
shortest_word_host	5.853749e-12
tld_in_path	1.049630e-12
nb_at	3.266972e-13
nb_and	5.624958e-15
domain_in_brand	2.764751e-17
nb_hyperlinks	4.243675e-18
shortening_service	1.018435e-19
nb_eq	8.097056e-22
links_in_tags	8.398174e-24
ratio_extMedia	1.147507e-24
external_favicon	6.272767e-26
suspicious_tld	1.456504e-26
safe_anchor	2.656364e-27
ratio_intHyperlinks	3.627670e-29
dns_record	3.751270e-30
abnormal_subdomain	1.708552e-30
domain_with_copyright	1.050737e-35
ratio_intMedia	4.054317e-39
statistical_report	1.245513e-39
domain_age	2.171043e-42
ratio_digits_host	3.005163e-44

phish_hints	6.311949e-50
domain_in_title	2.494145e-54
nb_qm	4.966874e-57
ratio_digits_url	1.363905e-59
empty_title	1.981960e-72
prefix_suffix	4.477458e-74
tld_in_subdomain	3.742944e-80
page_rank	1.484835e-105
nb_www	4.429634e-111
ip	2.154560e-174
google_index	0.000000e+00
nb_or	NaN
ratio_nullHyperlinks	NaN
ratio_intRedirection	NaN
ratio_intErrors	NaN
submit_email	NaN
sfh	NaN
dtype:	float64

```
[ ]: def DropFeature (p_values, threshold):
    drop_feature = set()
    for index, values in p_values.items():
        if values > threshold or np.isnan(values):
            drop_feature.add(index)
    return drop_feature
```

```
[ ]: drop_feature = DropFeature(p_values, .05)
len(set(drop_feature))
```

```
[ ]: 23
```

```
[ ]: drop_feature
```

```
[ ]: {'char_repeat',
      'iframe',
      'login_form',
      'nb_comma',
      'nb_dollar',
      'nb_or',
      'nb_percent',
      'nb_redirection',
      'nb_space',
      'nb_underscore',
      'onmouseover',
      'path_extension',
      'port',
      'punycode',
```



```

'random_domain',
'ratio_extErrors',
'ratio_intErrors',
'ratio_intRedirection',
'ratio_nullHyperlinks',
'right_clic',
'sfh',
'shortest_words_raw',
'submit_email'}

```

```

[ ]: X_train.drop(drop_feature, axis=1, inplace=True)
X_test.drop(drop_feature, axis=1, inplace=True)

```

```

[ ]: len(X_train.columns)

```

```

[ ]: 64

```

```

[ ]: len(X_test.columns)

```

```

[ ]: 64

```

```

[ ]: print("Training set has {} samples.".format(X_train.shape[0]))
print("Testing set has {} samples.".format(X_test.shape[0]))

```

Training set has 9144 samples.

Testing set has 2286 samples.

```

[ ]: from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression

# defining parameter range
param_grid = {'penalty' : ['l2'],
              'C' : [0.1, 1, 10, 20, 30],
              'solver' : ['lbfgs', 'newton-cg', 'liblinear', 'sag', 'saga'],
              'max_iter' : [2500, 5000]}

grid_logr = GridSearchCV(LogisticRegression(), param_grid, refit = True, cv = 10, verbose = 3, n_jobs = -1)

# fitting the model for grid search
grid_logr.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_logr.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_logr.best_estimator_)
print(grid_logr.best_score_)

```

```
Fitting 10 folds for each of 50 candidates, totalling 500 fits
{'C': 30, 'max_iter': 2500, 'penalty': 'l2', 'solver': 'saga'}
LogisticRegression(C=30, max_iter=2500, solver='saga')
0.9425849266420346
```

```
[ ]: log_r_model = grid_logr.best_estimator_

# Performing training
#log_r_model = log_r.fit(X_train, y_train.values.ravel())
```

```
[ ]: log_r_predict = log_r_model.predict(X_test)
```

```
[ ]: # from sklearn.metrics import confusion_matrix, accuracy_score
# cm = confusion_matrix(y_test, log_r_predict)
# ac = accuracy_score(y_test, log_r_predict)
```


```
[ ]: print ("Accuracy of log_r classifier : ", accuracy_score(y_test, log_r_predict)*100)
```

Accuracy of log\_r classifier : 94.35695538057742

```
[ ]: print(classification_report(y_test, log_r_predict))
```

	precision	recall	f1-score	support
0	0.94	0.95	0.94	1152
1	0.95	0.94	0.94	1134
accuracy			0.94	2286
macro avg	0.94	0.94	0.94	2286
weighted avg	0.94	0.94	0.94	2286

```
[ ]: sns.heatmap(confusion_matrix(y_test, log_r_predict), annot=True, fmt='g', cmap='Blues')
plt.title("LogisticRegression")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
```



```
[ ]: # from sklearn.neighbors import KNeighborsClassifier

# #training_accuracy=[]
# test_accuracy=[]

# neighbors=range(1,10)
# ##values.ravel() converts vector y to flattened array
# for i in neighbors:
#     knn=KNeighborsClassifier(n_neighbors=i)
#     knn_model = knn.fit(X_train,y_train.values.ravel())
#     #training_accuracy.append(knn.score(X_train,y_train.values.ravel()))
#     test_accuracy.append(knn_model.score(X_test,y_test.values.ravel()))
```

```
[ ]: # plt.plot(neighbors,test_accuracy,label="test accuracy")
# plt.ylabel("Accuracy")
# plt.xlabel("number of neighbors")
# plt.legend()
# plt.show()
```

```
[ ]: from sklearn.neighbors import KNeighborsClassifier

# defining parameter range
param_grid = {'n_neighbors': [1,2,3,4,5,6,7,8,9,10]}

grid_knn = GridSearchCV(KNeighborsClassifier(), param_grid, refit = True, cv = 10, verbose = 3, n_jobs = -1)

# fitting the model for grid search
grid_knn.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_knn.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_knn.best_estimator_)
print(grid_knn.best_score_)
```

```
Fitting 10 folds for each of 10 candidates, totalling 100 fits
{'n_neighbors': 3}
KNeighborsClassifier(n_neighbors=3)
0.9231213306070714
```

```
[ ]: knn_model = grid_knn.best_estimator_
#knn_model = knn.fit(X_train,y_train.values.ravel())
```

```
[ ]: #print ("Accuracy of knn classifier: ", max(test_accuracy)*100)
knn_predict = knn_model.predict(X_test)
```

```
[ ]: print('The accuracy of knn Classifier is: ', 100.0 * accuracy_score(y_test, knn_predict))
```


```
The accuracy of knn Classifier is: 92.38845144356955
```

```
[ ]: print(classification_report(y_test, knn_predict))
```

	precision	recall	f1-score	support
0	0.91	0.94	0.93	1152
1	0.94	0.91	0.92	1134
accuracy			0.92	2286
macro avg	0.92	0.92	0.92	2286
weighted avg	0.92	0.92	0.92	2286

```
[ ]: sns.heatmap(confusion_matrix(y_test, knn_predict), annot=True, fmt='g', cmap='Blues')
plt.title("KNeighborsClassifier")
```

```
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
```



```
[ ]: # # here is the change
# knn_y_pred_proba = knn.predict_proba(X_test)
# knn_y_pred_proba_positive = knn_y_pred_proba[:, 1]

# RocCurveDisplay.from_predictions(y_test,knn_y_pred_proba_positive)

# fig, ax = plt.subplots()
# RocCurveDisplay.from_estimator(
#     logreg, X_test, y_test, ax = ax)

# logreg_y_decision = logreg.decision_function(X_test)
# metrics.RocCurveDisplay.
↪from_predictions(y_test,logreg_y_decision,ax=ax,name="logreg predictions")
```

```
[ ]: from sklearn.svm import SVC

# defining parameter range
param_grid = {'C': [0.1, 1, 10],
              'gamma': [1, 0.1, 0.01],
              'kernel': ['linear', 'poly', 'rbf', 'sigmoid']}

grid_svc = GridSearchCV(SVC(), param_grid, refit = True, cv = 10, verbose = 3,
                        ↪n_jobs = -1)

# fitting the model for grid search
grid_svc.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_svc.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_svc.best_estimator_)
print(grid_svc.best_score_)
```

Fitting 10 folds for each of 36 candidates, totalling 360 fits  
{'C': 10, 'gamma': 0.1, 'kernel': 'rbf'}  
SVC(C=10, gamma=0.1)  
0.9572385837787423

```
[ ]: svc_model = grid_svc.best_estimator_
#svc_model = svc.fit(X_train,y_train.values.ravel())
```

```
[ ]: svc_predict = svc_model.predict(X_test)
```


```
[ ]: print('The accuracy of svc Classifier is: ', 100.0 * accuracy_score(y_test,
↪svc_predict))
```

The accuracy of svc Classifier is: 95.71303587051618

```
[ ]: print(classification_report(y_test, svc_predict))
```

	precision	recall	f1-score	support
0	0.95	0.97	0.96	1152
1	0.97	0.95	0.96	1134
accuracy			0.96	2286
macro avg	0.96	0.96	0.96	2286
weighted avg	0.96	0.96	0.96	2286

```
[ ]: sns.heatmap(confusion_matrix(y_test, svc_predict), annot=True, fmt='g',
    ↪ cmap='Blues')
plt.title("SVC")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
```



```
[ ]: from sklearn.svm import NuSVC

# defining parameter range
param_grid = {'nu': [0.1, 0.5],
              'gamma': [1, 0.1, 0.01],
              'kernel': ['rbf']} # 'linear', 'poly', 'rbf', 'sigmoid'

grid_nusvc = GridSearchCV(NuSVC(), param_grid, refit = True, verbose = 3, cv =
    ↪ 10, n_jobs = -1)

# fitting the model for grid search
grid_nusvc.fit(X_train, y_train.values.ravel())
```

```
# print best parameter after tuning
print(grid_nusvc.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_nusvc.best_estimator_)
print(grid_nusvc.best_score_)
```

```
Fitting 10 folds for each of 6 candidates, totalling 60 fits
{'gamma': 0.1, 'kernel': 'rbf', 'nu': 0.1}
NuSVC(gamma=0.1, nu=0.1)
0.9580039698198037
```

```
[ ]: nusvc_model = grid_nusvc.best_estimator_
#nusvc_model = nusvc.fit(X_train, y_train.values.ravel())
```

```
[ ]: nusvc_predict = nusvc_model.predict(X_test)
```

```
[ ]: print('The accuracy of nusvc Classifier is: ', 100.0 * accuracy_score(y_test,
↪nusvc_predict))
```


```
The accuracy of nusvc Classifier is: 95.71303587051618
```

```
[ ]: print(classification_report(y_test, nusvc_predict))
```

	precision	recall	f1-score	support
0	0.95	0.97	0.96	1152
1	0.96	0.95	0.96	1134
accuracy			0.96	2286
macro avg	0.96	0.96	0.96	2286
weighted avg	0.96	0.96	0.96	2286

```
[ ]: sns.heatmap(confusion_matrix(y_test, nusvc_predict), annot=True, fmt='g',
↪cmap='Blues')
plt.title("NuSVC")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
```





```
[ ]: from sklearn.svm import LinearSVC

# defining parameter range
param_grid = {'C': [0.1, 1, 10, 20, 30],
              'penalty': ['l1', 'l2'],
              'loss': ['squared_hinge'],
              'dual': [False],
              'tol': [.1, .01, .001]}

grid_lsvc = GridSearchCV(LinearSVC(), param_grid, refit = True, verbose = 3, cv=
↳ 10, n_jobs = -1)

# fitting the model for grid search
grid_lsvc.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_lsvc.best_params_)

# print how our model looks after hyper-parameter tuning
```

```
print(grid_lsvc.best_estimator_)
print(grid_lsvc.best_score_)
```

Fitting 10 folds for each of 30 candidates, totalling 300 fits  
 {'C': 20, 'dual': False, 'loss': 'squared\_hinge', 'penalty': 'l2', 'tol': 0.001}  
 LinearSVC(C=20, dual=False, tol=0.001)  
 0.9426942162595209

```
[ ]: lsvc_model = grid_lsvc.best_estimator_
      #lsvc_model = lsvc.fit(X_train, y_train.values.ravel())
```

```
[ ]: lsvc_predict = lsvc_model.predict(X_test)
```

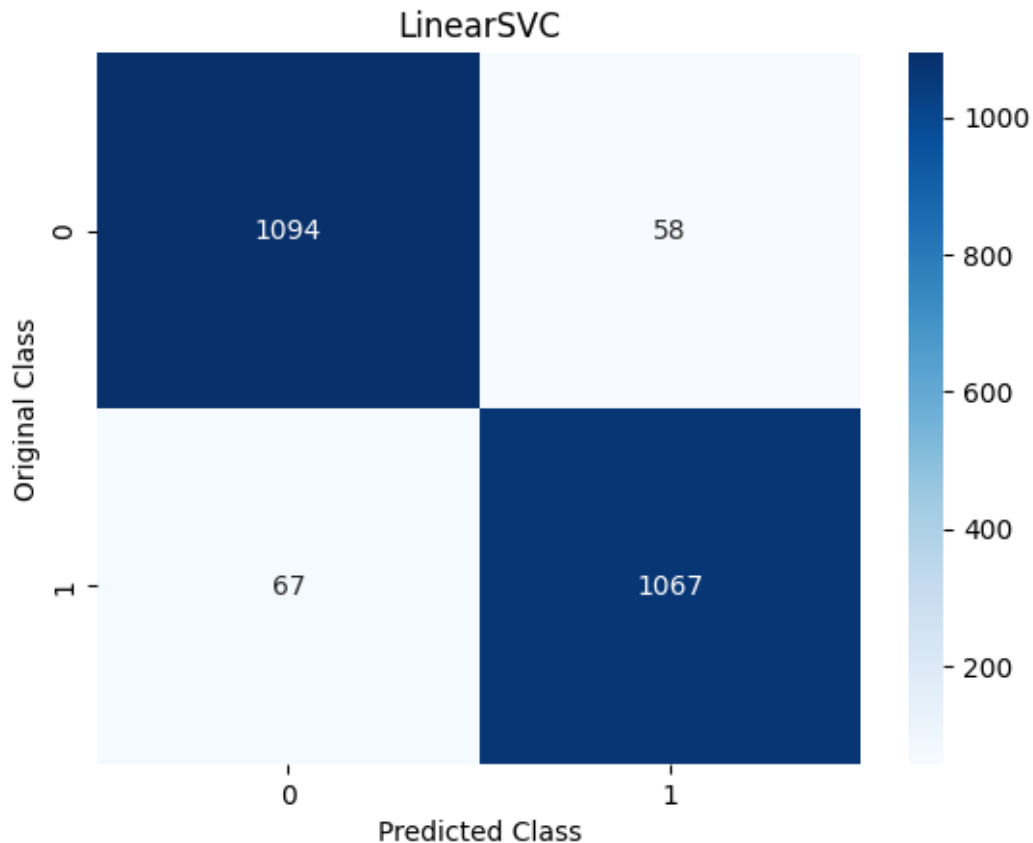
```
[ ]: print('The accuracy of lsvc Classifier is: ', 100.0 * accuracy_score(y_test,
      ↪lsvc_predict))
```

The accuracy of lsvc Classifier is: 94.53193350831145

```
[ ]: print(classification_report(y_test, lsvc_predict))
```

	precision	recall	f1-score	support
0	0.94	0.95	0.95	1152
1	0.95	0.94	0.94	1134
accuracy			0.95	2286
macro avg	0.95	0.95	0.95	2286
weighted avg	0.95	0.95	0.95	2286

```
[ ]: sns.heatmap(confusion_matrix(y_test, lsvc_predict), annot=True, fmt='g',
      ↪cmap='Blues')
      plt.title("LinearSVC")
      plt.xlabel('Predicted Class')
      plt.ylabel('Original Class')
      plt.show()
```



```
[ ]: from sklearn.ensemble import AdaBoostClassifier

# defining parameter range
param_grid = {'n_estimators': [40,50,100,200,300]}

grid_ada = GridSearchCV(AdaBoostClassifier(), param_grid, refit = True, verbose=
    ↪ 3, cv = 10, n_jobs = -1)

# fitting the model for grid search
grid_ada.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_ada.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_ada.best_estimator_)
print(grid_ada.best_score_)
```

Fitting 10 folds for each of 5 candidates, totalling 50 fits  
 {'n\_estimators': 300}

```
AdaBoostClassifier(n_estimators=300)
0.9539591778168383
```

```
[ ]: ada_model = grid_ada.best_estimator_  
      #ada_model = ada.fit(X_train,y_train.values.ravel())
```

```
[ ]: ada_predict = ada_model.predict(X_test)
```


```
[ ]: print('The accuracy of Ada Boost Classifier is: ', 100.0 *  
      ↪accuracy_score(ada_predict,y_test))
```

The accuracy of Ada Boost Classifier is: 95.2755905511811

```
[ ]: print(classification_report(y_test, ada_predict))
```

	precision	recall	f1-score	support
0	0.95	0.96	0.95	1152
1	0.96	0.95	0.95	1134
accuracy			0.95	2286
macro avg	0.95	0.95	0.95	2286
weighted avg	0.95	0.95	0.95	2286

```
[ ]: sns.heatmap(confusion_matrix(y_test, ada_predict), annot=True, fmt='g',  
      ↪cmap='Blues')  
plt.title("AdaBoostClassifier")  
plt.xlabel('Predicted Class')  
plt.ylabel('Original Class')  
plt.show()
```



```
[ ]: from xgboost import XGBClassifier

# defining parameter range
param_grid = {
    "gamma": [.01, .1, .5],
    "n_estimators": [50,100,150,200,250]
}

grid_xgb = GridSearchCV(XGBClassifier(), param_grid, refit = True, verbose = 3,
    ↪cv = 10, n_jobs = -1)

# fitting the model for grid search
grid_xgb.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_xgb.best_params_)

# print how our model looks after hyper-parameter tuning
```

```
print(grid_xgb.best_estimator_)
print(grid_xgb.best_score_)
```

Fitting 10 folds for each of 15 candidates, totalling 150 fits

```
{'gamma': 0.1, 'n_estimators': 150}
```

```
XGBClassifier(base_score=0.5, booster='gbtree', callbacks=None,
              colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,
              early_stopping_rounds=None, enable_categorical=False,
              eval_metric=None, gamma=0.1, gpu_id=-1, grow_policy='depthwise',
              importance_type=None, interaction_constraints='',
              learning_rate=0.300000012, max_bin=256, max_cat_to_onehot=4,
              max_delta_step=0, max_depth=6, max_leaves=0, min_child_weight=1,
              missing=nan, monotone_constraints='()', n_estimators=150,
              n_jobs=0, num_parallel_tree=1, predictor='auto', random_state=0,
              reg_alpha=0, reg_lambda=1, ...)
```

0.9704711171694707

```
[ ]: xgb_model = grid_xgb.best_estimator_
      #xgb_model = xgb.fit(X_train,y_train)
```

```
[ ]: xgb_predict=xgb_model.predict(X_test)
```


```
[ ]: print('The accuracy of XGBoost Classifier is: ', 100.0 *
      ↪accuracy_score(xgb_predict,y_test))
```

The accuracy of XGBoost Classifier is: 96.93788276465442

```
[ ]: print(classification_report(y_test, xgb_predict))
```

	precision	recall	f1-score	support
0	0.96	0.97	0.97	1152
1	0.97	0.96	0.97	1134
accuracy			0.97	2286
macro avg	0.97	0.97	0.97	2286
weighted avg	0.97	0.97	0.97	2286

```
[ ]: sns.heatmap(confusion_matrix(y_test, xgb_predict), annot=True, fmt='g',
      ↪cmap='Blues')
plt.title("XGBClassifier")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
```



```
[ ]: from sklearn.ensemble import GradientBoostingClassifier

# defining parameter range
param_grid = {
    "learning_rate": [.1,.5,1],
    "n_estimators": [50,100,150,200,250]
}

grid_gbc = GridSearchCV(GradientBoostingClassifier(), param_grid, refit = True,
    verbose = 3, cv = 10, n_jobs = -1)

# fitting the model for grid search
grid_gbc.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_gbc.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_gbc.best_estimator_)
```

```
print(grid_gbc.best_score_)
```

Fitting 10 folds for each of 15 candidates, totalling 150 fits  
{'learning\_rate': 0.5, 'n\_estimators': 250}  
GradientBoostingClassifier(learning\_rate=0.5, n\_estimators=250)  
0.9658795183604166

```
[ ]: gbc_model = grid_gbc.best_estimator_  
      #gbc_model = gbc.fit(X_train,y_train.values.ravel())  
  
      #clf = GradientBoostingClassifier(n_estimators=100, learning_rate=1.0,  
      #    max_depth=1, random_state=0).fit(X_train, y_train)  
      #clf.score(X_test, y_test)
```

```
[ ]: gbc_predict = gbc_model.predict(X_test)
```

```
[ ]: print('The accuracy of GradientBoost Classifier is: ', 100.0 *  
      ↪accuracy_score(gbc_predict,y_test))
```


The accuracy of GradientBoost Classifier is: 96.63167104111986

```
[ ]: print(classification_report(y_test, gbc_predict))
```

	precision	recall	f1-score	support
0	0.96	0.97	0.97	1152
1	0.97	0.96	0.97	1134
accuracy			0.97	2286
macro avg	0.97	0.97	0.97	2286
weighted avg	0.97	0.97	0.97	2286

```
[ ]: sns.heatmap(confusion_matrix(y_test, gbc_predict), annot=True, fmt='g',  
      ↪cmap='Blues')  
plt.title("GradientBoostingClassifier")  
plt.xlabel('Predicted Class')  
plt.ylabel('Original Class')  
plt.show()
```





```
[ ]: # gbc_model.get_params().keys()
```

```
[ ]: # import inspect
# import sklearn
# import xgboost

# models = [xgboost.XGBClassifier]
# for m in models:
#     hyperparams = inspect.signature(m.__init__)
#     print(hyperparams)
# #or
# xgb_model.get_params().keys()
```

```
[ ]: from sklearn.ensemble import BaggingClassifier
from sklearn.tree import DecisionTreeClassifier

# defining parameter range
param_grid = {
    "base_estimator": [DecisionTreeClassifier()],
    "n_estimators": [50,100,150,200,250]
```

```

}

grid_bag = GridSearchCV(BaggingClassifier(), param_grid, refit = True, verbose_
↳ = 3, cv = 10, n_jobs = -1)

# fitting the model for grid search
grid_bag.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_bag.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_bag.best_estimator_)
print(grid_bag.best_score_)

```

```

Fitting 10 folds for each of 5 candidates, totalling 50 fits
{'base_estimator': DecisionTreeClassifier(), 'n_estimators': 250}
BaggingClassifier(base_estimator=DecisionTreeClassifier(), n_estimators=250)
0.9575681266516005

```

```

[ ]: bag_model = grid_bag.best_estimator_
      #bag_model = bag.fit(X_train, y_train.values.ravel())

```

```

[ ]: bag_predict = bag_model.predict(X_test)

```

```

[ ]: print('The accuracy of Bagging Classifier is: ', 100.0 * _
↳ accuracy_score(y_test, bag_predict))

```

```

The accuracy of Bagging Classifier is: 95.84426946631672

```

```

[ ]: print(classification_report(y_test, bag_predict))


```

	precision	recall	f1-score	support
0	0.95	0.96	0.96	1152
1	0.96	0.95	0.96	1134
accuracy			0.96	2286
macro avg	0.96	0.96	0.96	2286
weighted avg	0.96	0.96	0.96	2286

```

[ ]: sns.heatmap(confusion_matrix(y_test, bag_predict), annot=True, fmt='g', _
↳ cmap='Blues')
plt.title("BaggingClassifier")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()

```



```
[ ]: from sklearn.ensemble import RandomForestClassifier

# defining parameter range
param_grid = {
    "n_estimators": [50,100,150,200,250]
}

grid_rfc = GridSearchCV(RandomForestClassifier(), param_grid, refit = True,
    verbose = 3, cv = 10, n_jobs = -1)

# fitting the model for grid search
grid_rfc.fit(X_train, y_train.values.ravel())

# print best parameter after tuning
print(grid_rfc.best_params_)

# print how our model looks after hyper-parameter tuning
print(grid_rfc.best_estimator_)
print(grid_rfc.best_score_)
```

```
Fitting 10 folds for each of 5 candidates, totalling 50 fits
{'n_estimators': 100}
RandomForestClassifier()
0.9666444261099352
```

```
[ ]: rfc_model = grid_rfc.best_estimator_
      #rfc_model = rfc.fit(X_train,y_train.values.ravel())
```

```
[ ]: rfc_predict = rfc_model.predict(X_test)
```


```
[ ]: print('The accuracy of RandomForest Classifier is: ' , 100.0 *
      ↪accuracy_score(rfc_predict,y_test))
```

The accuracy of RandomForest Classifier is: 96.23797025371829


```
[ ]: print(classification_report(y_test, rfc_predict))
```

	precision	recall	f1-score	support
0	0.96	0.97	0.96	1152
1	0.97	0.96	0.96	1134
accuracy			0.96	2286
macro avg	0.96	0.96	0.96	2286
weighted avg	0.96	0.96	0.96	2286

```
[ ]: sns.heatmap(confusion_matrix(y_test, rfc_predict), annot=True, fmt='g',
      ↪cmap='Blues')
plt.title("RandomForestClassifier")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
```



```
[ ]: estimators =  
    ↳ [logr_model, knn_model, svc_model, nusvc_model, lsvc_model, xgb_model, ada_model, gbc_model, bag_mo  
  
for estimator in estimators:  
    RocCurveDisplay.from_estimator(estimator, X_test, y_test, ax=plt.gca())
```



```
[ ]: import tensorflow as tf
      #from tensorflow.keras.datasets import imdb
      from keras.layers import Embedding, Dense, LSTM, BatchNormalization
      from keras.losses import BinaryCrossentropy
      from keras.models import Sequential
      from keras.optimizers import Adam
      #from tensorflow.keras.preprocessing.sequence import pad_sequences

      # Model configuration
      additional_metrics = ['accuracy']
      batch_size = 32
      #embedding_output_dims = (X_train.shape[1])
      loss_function = BinaryCrossentropy()
      #max_sequence_length = (X_train.shape[1])
      #num_distinct_words = (X_train.shape[1])
      number_of_epochs = 100
      optimizer = Adam()
      validation_split = 0.20
      verbosity_mode = 1

      # reshape from [samples, features] into [samples, timesteps, features]
```

```

timesteps = 1
X_train_reshape = X_train.values.ravel().reshape(X_train.shape[0],timesteps,
↳X_train.shape[1])
X_test_reshape = X_test.values.ravel().reshape(X_test.shape[0],timesteps,
↳X_test.shape[1])

# Disable eager execution
#tf.compat.v1.disable_eager_execution()

# Load dataset
# (x_train, y_train), (x_test, y_test) = imdb.
↳load_data(num_words=num_distinct_words)
# print(x_train.shape)
# print(x_test.shape)

# Pad all sequences
# padded_inputs = pad_sequences(X_train, maxlen=max_sequence_length, value = 0.
↳0) # 0.0 because it corresponds with <PAD>
# padded_inputs_test = pad_sequences(X_test, maxlen=max_sequence_length, value
↳= 0.0) # 0.0 because it corresponds with <PAD>

# Define the Keras model
def build_model_lstm():
    model = Sequential()
    #model.add(Embedding(num_distinct_words, embedding_output_dims,
↳input_length=max_sequence_length))
    model.add(LSTM(100, input_shape = (timesteps,X_train_reshape.shape[2])))
    model.add(BatchNormalization())
    model.add(Dense(50, activation='relu'))
    model.add(Dense(25, activation='relu'))
    model.add(Dense(10, activation='relu'))
    model.add(Dense(1, activation='sigmoid'))

    # Compile the model
    model.compile(optimizer=optimizer, loss=loss_function,
↳metrics=additional_metrics)
    return model

#from keras.wrappers.scikit_learn import KerasClassifier
lstm_model = build_model_lstm()
# Give a summary
lstm_model.summary()

# Train the model

```

```

history = lstm_model.fit(X_train_reshape, y_train.values.ravel(),
    ↪batch_size=batch_size, epochs=number_of_epochs, verbose=verbosity_mode,
    ↪validation_split=validation_split)

# Test the model after training
#lstm_predict = lstm_model.predict(X_test_reshape)
test_results = lstm_model.evaluate(X_test_reshape, y_test.values.ravel(),
    ↪verbose=False)
print(f'Test results - Loss: {test_results[0]} - Accuracy:
    ↪{100*test_results[1]}%')

```

Model: "sequential"

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, 100)	66000
batch_normalization (Batch Normalization)	(None, 100)	400
dense (Dense)	(None, 50)	5050

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, 100)	66000
batch_normalization (Batch Normalization)	(None, 100)	400
dense (Dense)	(None, 50)	5050
dense_1 (Dense)	(None, 25)	1275
dense_2 (Dense)	(None, 10)	260
dense_3 (Dense)	(None, 1)	11

Total params: 72,996  
 Trainable params: 72,796  
 Non-trainable params: 200

Epoch 1/100  
 229/229 [=====] - 4s 7ms/step - loss: 0.2223 -  
 accuracy: 0.9136 - val\_loss: 0.4193 - val\_accuracy: 0.9295  
 Epoch 2/100



229/229 [=====] - 1s 4ms/step - loss: 0.1568 -  
accuracy: 0.9423 - val\_loss: 0.2086 - val\_accuracy: 0.9399  
Epoch 3/100  
229/229 [=====] - 1s 4ms/step - loss: 0.1376 -  
accuracy: 0.9489 - val\_loss: 0.1596 - val\_accuracy: 0.9404  
Epoch 4/100  
229/229 [=====] - 1s 4ms/step - loss: 0.1270 -  
accuracy: 0.9535 - val\_loss: 0.1526 - val\_accuracy: 0.9442  
Epoch 5/100  
229/229 [=====] - 1s 4ms/step - loss: 0.1111 -  
accuracy: 0.9571 - val\_loss: 0.1533 - val\_accuracy: 0.9459  
Epoch 6/100  
229/229 [=====] - 1s 4ms/step - loss: 0.1045 -  
accuracy: 0.9610 - val\_loss: 0.1502 - val\_accuracy: 0.9497  
Epoch 7/100  
229/229 [=====] - 1s 4ms/step - loss: 0.1006 -  
accuracy: 0.9628 - val\_loss: 0.1514 - val\_accuracy: 0.9492  
Epoch 8/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0983 -  
accuracy: 0.9613 - val\_loss: 0.1578 - val\_accuracy: 0.9399  
Epoch 9/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0953 -  
accuracy: 0.9632 - val\_loss: 0.1507 - val\_accuracy: 0.9486  
Epoch 10/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0871 -  
accuracy: 0.9668 - val\_loss: 0.1687 - val\_accuracy: 0.9442  
Epoch 11/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0802 -  
accuracy: 0.9714 - val\_loss: 0.1812 - val\_accuracy: 0.9426  
Epoch 12/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0722 -  
accuracy: 0.9729 - val\_loss: 0.1630 - val\_accuracy: 0.9541  
Epoch 13/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0682 -  
accuracy: 0.9733 - val\_loss: 0.1557 - val\_accuracy: 0.9535  
Epoch 14/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0654 -  
accuracy: 0.9754 - val\_loss: 0.1801 - val\_accuracy: 0.9486  
Epoch 15/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0701 -  
accuracy: 0.9721 - val\_loss: 0.1736 - val\_accuracy: 0.9502  
Epoch 16/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0623 -  
accuracy: 0.9765 - val\_loss: 0.1818 - val\_accuracy: 0.9546  
Epoch 17/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0597 -  
accuracy: 0.9761 - val\_loss: 0.1913 - val\_accuracy: 0.9453  
Epoch 18/100

229/229 [=====] - 1s 4ms/step - loss: 0.0554 -  
accuracy: 0.9813 - val\_loss: 0.1735 - val\_accuracy: 0.9502  
Epoch 19/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0521 -  
accuracy: 0.9821 - val\_loss: 0.1722 - val\_accuracy: 0.9524  
Epoch 20/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0537 -  
accuracy: 0.9796 - val\_loss: 0.1708 - val\_accuracy: 0.9459  
Epoch 21/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0496 -  
accuracy: 0.9824 - val\_loss: 0.1913 - val\_accuracy: 0.9431  
Epoch 22/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0508 -  
accuracy: 0.9795 - val\_loss: 0.1987 - val\_accuracy: 0.9431  
Epoch 23/100  
229/229 [=====] - 1s 5ms/step - loss: 0.0451 -  
accuracy: 0.9828 - val\_loss: 0.1873 - val\_accuracy: 0.9475  
Epoch 24/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0455 -  
accuracy: 0.9829 - val\_loss: 0.1895 - val\_accuracy: 0.9519  
Epoch 25/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0409 -  
accuracy: 0.9846 - val\_loss: 0.1957 - val\_accuracy: 0.9519  
Epoch 26/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0411 -  
accuracy: 0.9862 - val\_loss: 0.2121 - val\_accuracy: 0.9453  
Epoch 27/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0376 -  
accuracy: 0.9870 - val\_loss: 0.2127 - val\_accuracy: 0.9502  
Epoch 28/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0349 -  
accuracy: 0.9882 - val\_loss: 0.2015 - val\_accuracy: 0.9453  
Epoch 29/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0425 -  
accuracy: 0.9848 - val\_loss: 0.1991 - val\_accuracy: 0.9481  
Epoch 30/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0362 -  
accuracy: 0.9877 - val\_loss: 0.2377 - val\_accuracy: 0.9470  
Epoch 31/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0364 -  
accuracy: 0.9867 - val\_loss: 0.2108 - val\_accuracy: 0.9486  
Epoch 32/100  
229/229 [=====] - 1s 5ms/step - loss: 0.0343 -  
accuracy: 0.9866 - val\_loss: 0.2275 - val\_accuracy: 0.9442  
Epoch 33/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0292 -  
accuracy: 0.9892 - val\_loss: 0.2329 - val\_accuracy: 0.9502  
Epoch 34/100

229/229 [=====] - 1s 4ms/step - loss: 0.0330 -  
accuracy: 0.9888 - val\_loss: 0.2292 - val\_accuracy: 0.9453  
Epoch 35/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0289 -  
accuracy: 0.9897 - val\_loss: 0.2268 - val\_accuracy: 0.9530  
Epoch 36/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0338 -  
accuracy: 0.9869 - val\_loss: 0.2525 - val\_accuracy: 0.9371  
Epoch 37/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0275 -  
accuracy: 0.9906 - val\_loss: 0.2191 - val\_accuracy: 0.9519  
Epoch 38/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0295 -  
accuracy: 0.9889 - val\_loss: 0.2277 - val\_accuracy: 0.9442  
Epoch 39/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0217 -  
accuracy: 0.9928 - val\_loss: 0.2287 - val\_accuracy: 0.9535  
Epoch 40/100  
229/229 [=====] - 1s 5ms/step - loss: 0.0210 -  
accuracy: 0.9933 - val\_loss: 0.2635 - val\_accuracy: 0.9431  
Epoch 41/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0240 -  
accuracy: 0.9910 - val\_loss: 0.2428 - val\_accuracy: 0.9481  
Epoch 42/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0271 -  
accuracy: 0.9907 - val\_loss: 0.2515 - val\_accuracy: 0.9568  
Epoch 43/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0245 -  
accuracy: 0.9921 - val\_loss: 0.2799 - val\_accuracy: 0.9366  
Epoch 44/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0201 -  
accuracy: 0.9926 - val\_loss: 0.2768 - val\_accuracy: 0.9475  
Epoch 45/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0235 -  
accuracy: 0.9915 - val\_loss: 0.2470 - val\_accuracy: 0.9502  
Epoch 46/100  
229/229 [=====] - 1s 5ms/step - loss: 0.0233 -  
accuracy: 0.9906 - val\_loss: 0.2689 - val\_accuracy: 0.9475  
Epoch 47/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0214 -  
accuracy: 0.9919 - val\_loss: 0.2691 - val\_accuracy: 0.9481  
Epoch 48/100  
229/229 [=====] - 1s 5ms/step - loss: 0.0208 -  
accuracy: 0.9922 - val\_loss: 0.2809 - val\_accuracy: 0.9541  
Epoch 49/100  
229/229 [=====] - 1s 4ms/step - loss: 0.0198 -  
accuracy: 0.9930 - val\_loss: 0.2851 - val\_accuracy: 0.9486  
Epoch 50/100

229/229 [=====] - 1s 4ms/step - loss: 0.0208 -  
 accuracy: 0.9930 - val\_loss: 0.2764 - val\_accuracy: 0.9442  
 Epoch 51/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0206 -  
 accuracy: 0.9921 - val\_loss: 0.2710 - val\_accuracy: 0.9453  
 Epoch 52/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0237 -  
 accuracy: 0.9922 - val\_loss: 0.2615 - val\_accuracy: 0.9519  
 Epoch 53/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0189 -  
 accuracy: 0.9937 - val\_loss: 0.2734 - val\_accuracy: 0.9513  
 Epoch 54/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0131 -  
 accuracy: 0.9960 - val\_loss: 0.3044 - val\_accuracy: 0.9426  
 Epoch 55/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0160 -  
 accuracy: 0.9940 - val\_loss: 0.2800 - val\_accuracy: 0.9519  
 Epoch 56/100  
 229/229 [=====] - 1s 4ms/step - loss: 0.0178 -  
 accuracy: 0.9929 - val\_loss: 0.3062 - val\_accuracy: 0.9393  
 Epoch 57/100  
 229/229 [=====] - 1s 4ms/step - loss: 0.0238 -  
 accuracy: 0.9914 - val\_loss: 0.2737 - val\_accuracy: 0.9508  
 Epoch 58/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0164 -  
 accuracy: 0.9940 - val\_loss: 0.3027 - val\_accuracy: 0.9513  
 Epoch 59/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0239 -  
 accuracy: 0.9922 - val\_loss: 0.2741 - val\_accuracy: 0.9524  
 Epoch 60/100  
 229/229 [=====] - 1s 4ms/step - loss: 0.0175 -  
 accuracy: 0.9941 - val\_loss: 0.2563 - val\_accuracy: 0.9524  
 Epoch 61/100  
 229/229 [=====] - 1s 4ms/step - loss: 0.0139 -  
 accuracy: 0.9944 - val\_loss: 0.2518 - val\_accuracy: 0.9513  
 Epoch 62/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0139 -  
 accuracy: 0.9955 - val\_loss: 0.2698 - val\_accuracy: 0.9492  
 Epoch 63/100  
 229/229 [=====] - 1s 4ms/step - loss: 0.0193 -  
 accuracy: 0.9937 - val\_loss: 0.2888 - val\_accuracy: 0.9420  
 Epoch 64/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0196 -  
 accuracy: 0.9930 - val\_loss: 0.3050 - val\_accuracy: 0.9475  
 Epoch 65/100  
 229/229 [=====] - 1s 4ms/step - loss: 0.0129 -  
 accuracy: 0.9952 - val\_loss: 0.2995 - val\_accuracy: 0.9470  
 Epoch 66/100

229/229 [=====] - 1s 5ms/step - loss: 0.0115 -  
 accuracy: 0.9962 - val\_loss: 0.2883 - val\_accuracy: 0.9470  
 Epoch 67/100  
 229/229 [=====] - 1s 4ms/step - loss: 0.0139 -  
 accuracy: 0.9945 - val\_loss: 0.3136 - val\_accuracy: 0.9497  
 Epoch 68/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0122 -  
 accuracy: 0.9955 - val\_loss: 0.3187 - val\_accuracy: 0.9486  
 Epoch 69/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0128 -  
 accuracy: 0.9944 - val\_loss: 0.3468 - val\_accuracy: 0.9481  
 Epoch 70/100  
 229/229 [=====] - 1s 4ms/step - loss: 0.0191 -  
 accuracy: 0.9928 - val\_loss: 0.3448 - val\_accuracy: 0.9431  
 Epoch 71/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0141 -  
 accuracy: 0.9952 - val\_loss: 0.3339 - val\_accuracy: 0.9420  
 Epoch 72/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0077 -  
 accuracy: 0.9973 - val\_loss: 0.3257 - val\_accuracy: 0.9502  
 Epoch 73/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0181 -  
 accuracy: 0.9937 - val\_loss: 0.3440 - val\_accuracy: 0.9431  
 Epoch 74/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0147 -  
 accuracy: 0.9945 - val\_loss: 0.3110 - val\_accuracy: 0.9459  
 Epoch 75/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0140 -  
 accuracy: 0.9954 - val\_loss: 0.3554 - val\_accuracy: 0.9437  
 Epoch 76/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0123 -  
 accuracy: 0.9959 - val\_loss: 0.3196 - val\_accuracy: 0.9497  
 Epoch 77/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0108 -  
 accuracy: 0.9962 - val\_loss: 0.3644 - val\_accuracy: 0.9464  
 Epoch 78/100  
 229/229 [=====] - 1s 4ms/step - loss: 0.0180 -  
 accuracy: 0.9941 - val\_loss: 0.3778 - val\_accuracy: 0.9420  
 Epoch 79/100  
 229/229 [=====] - 1s 4ms/step - loss: 0.0127 -  
 accuracy: 0.9958 - val\_loss: 0.3327 - val\_accuracy: 0.9420  
 Epoch 80/100  
 229/229 [=====] - 1s 4ms/step - loss: 0.0094 -  
 accuracy: 0.9967 - val\_loss: 0.3699 - val\_accuracy: 0.9475  
 Epoch 81/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0147 -  
 accuracy: 0.9943 - val\_loss: 0.3265 - val\_accuracy: 0.9486  
 Epoch 82/100

229/229 [=====] - 1s 5ms/step - loss: 0.0106 -  
 accuracy: 0.9963 - val\_loss: 0.3280 - val\_accuracy: 0.9481  
 Epoch 83/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0101 -  
 accuracy: 0.9967 - val\_loss: 0.3406 - val\_accuracy: 0.9486  
 Epoch 84/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0102 -  
 accuracy: 0.9967 - val\_loss: 0.3587 - val\_accuracy: 0.9464  
 Epoch 85/100  
 229/229 [=====] - 1s 4ms/step - loss: 0.0114 -  
 accuracy: 0.9958 - val\_loss: 0.3676 - val\_accuracy: 0.9442  
 Epoch 86/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0119 -  
 accuracy: 0.9952 - val\_loss: 0.3554 - val\_accuracy: 0.9410  
 Epoch 87/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0089 -  
 accuracy: 0.9974 - val\_loss: 0.3610 - val\_accuracy: 0.9492  
 Epoch 88/100  
 229/229 [=====] - 1s 4ms/step - loss: 0.0091 -  
 accuracy: 0.9970 - val\_loss: 0.4277 - val\_accuracy: 0.9404  
 Epoch 89/100  
 229/229 [=====] - 1s 4ms/step - loss: 0.0178 -  
 accuracy: 0.9937 - val\_loss: 0.3436 - val\_accuracy: 0.9508  
 Epoch 90/100  
 229/229 [=====] - 1s 4ms/step - loss: 0.0098 -  
 accuracy: 0.9967 - val\_loss: 0.3558 - val\_accuracy: 0.9513  
 Epoch 91/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0109 -  
 accuracy: 0.9962 - val\_loss: 0.4206 - val\_accuracy: 0.9415  
 Epoch 92/100  
 229/229 [=====] - 1s 4ms/step - loss: 0.0192 -  
 accuracy: 0.9941 - val\_loss: 0.3428 - val\_accuracy: 0.9453  
 Epoch 93/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0116 -  
 accuracy: 0.9962 - val\_loss: 0.3141 - val\_accuracy: 0.9492  
 Epoch 94/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0078 -  
 accuracy: 0.9970 - val\_loss: 0.3599 - val\_accuracy: 0.9459  
 Epoch 95/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0058 -  
 accuracy: 0.9978 - val\_loss: 0.3747 - val\_accuracy: 0.9475  
 Epoch 96/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0063 -  
 accuracy: 0.9971 - val\_loss: 0.3997 - val\_accuracy: 0.9459  
 Epoch 97/100  
 229/229 [=====] - 1s 5ms/step - loss: 0.0159 -  
 accuracy: 0.9958 - val\_loss: 0.3611 - val\_accuracy: 0.9420  
 Epoch 98/100

```

229/229 [=====] - 1s 5ms/step - loss: 0.0114 -
accuracy: 0.9958 - val_loss: 0.3565 - val_accuracy: 0.9464
Epoch 99/100
229/229 [=====] - 1s 4ms/step - loss: 0.0104 -
accuracy: 0.9959 - val_loss: 0.3956 - val_accuracy: 0.9464
Epoch 100/100
229/229 [=====] - 1s 4ms/step - loss: 0.0109 -
accuracy: 0.9966 - val_loss: 0.3779 - val_accuracy: 0.9502
Test results - Loss: 0.32264989614486694 - Accuracy: 95.31933665275574%

```

```

[ ]: lstm_predict_proba = lstm_model.predict(X_test_reshape, batch_size=32)
lstm_predict_class = (lstm_predict_proba > 0.5).astype("int32")
print(classification_report(y_test, lstm_predict_class))

```

```

72/72 [=====] - 1s 2ms/step


```

	precision	recall	f1-score	support
0	0.96	0.95	0.95	1152
1	0.95	0.96	0.95	1134
accuracy			0.95	2286
macro avg	0.95	0.95	0.95	2286
weighted avg	0.95	0.95	0.95	2286

```

[ ]: sns.heatmap(confusion_matrix(y_test, lstm_predict_class), annot=True, fmt='g',
cmap='Blues')
plt.title("LSTM")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()

```



```
[ ]: RocCurveDisplay.from_predictions(y_test,lstm_predict_class)
plt.show()
```



